

CSE 421/521 - Operating Systems
Fall 2014

LECTURE - II
OS STRUCTURES

Tevfik Koşar

University at Buffalo

August 28th, 2014

Roadmap

- Major OS Components
 - Processes and Threads
 - Memory management
 - CPU Scheduling
 - I/O Management
- Different OS Design Approaches
 - Simple Structure
 - Layered Approach
 - Microkernels
 - Modules



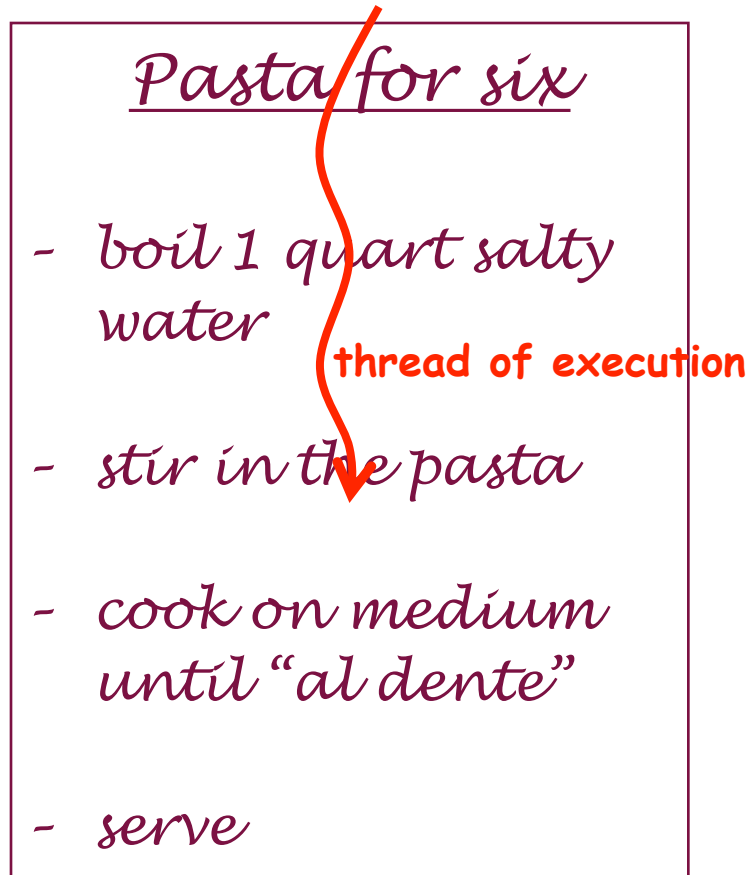
MAJOR OS COMPONENTS

Major OS Components


- Processes and Threads
- Memory management
- CPU Scheduling
- I/O Management

Processes and Threads

- a **Process** is a program in execution;



Program




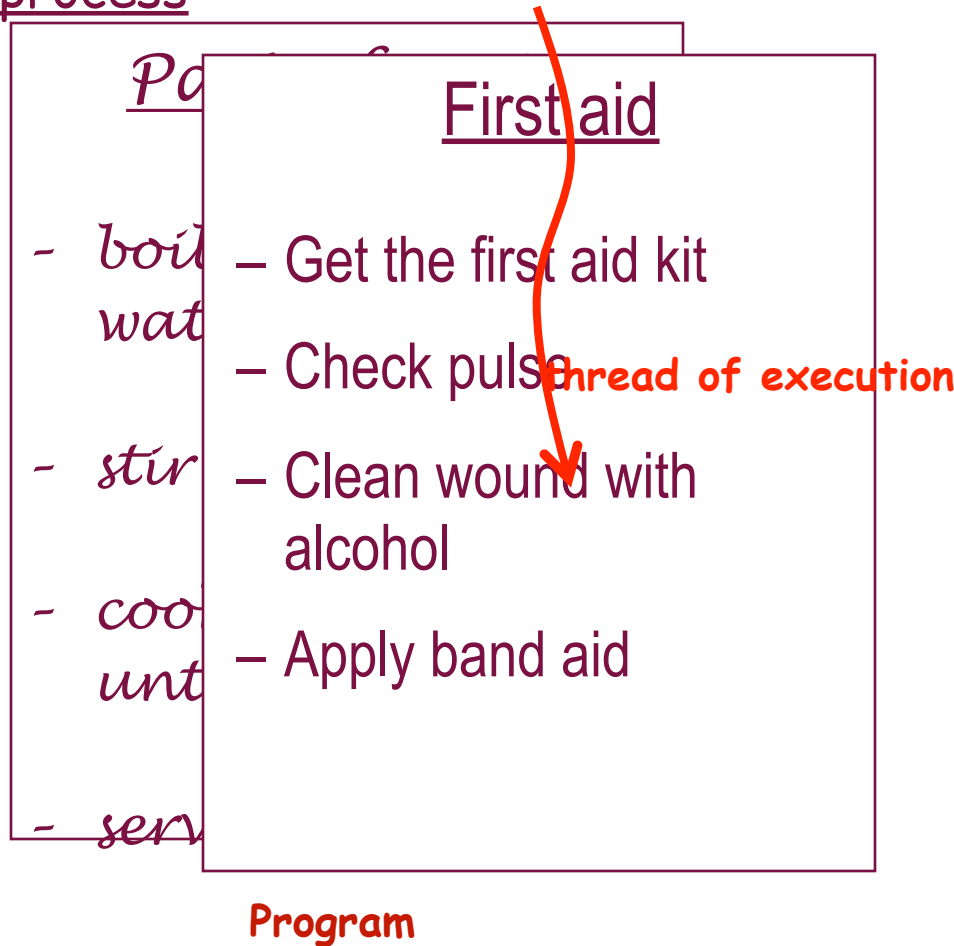
CPU

input data

Process

Processes and Threads

- It can be interrupted to let the CPU execute a higher-priority process



Processes and Threads


- ... and then resumed exactly where the CPU left off

Pasta for six

- boil 1 quart salty water
- stir in the pasta
- cook on medium until "al dente"
- serve


Program

thread of execution




Multitasking

- Multitasking gives the illusion of parallel processing (independent virtual program counters) on one CPU



(a) Multitasking from the CPU's viewpoint



(b) Multitasking from the processes' viewpoint = 4 virtual program counters

Pseudoparallelism in multitasking

TimeSharing

- **Timesharing** is logical extension in which CPU switches jobs so frequently that users can interact with each job while it is running, creating **interactive** computing
 - **Response time** should be < 1 second
 - Each user has at least one program loaded in memory and executing \Rightarrow **process**

Processes and Threads


➤ Operating System Responsibilities:

The O/S is responsible for managing processes and Threads

- ✓ the O/S creates & deletes processes and threads
- ✓ the O/S suspends & resumes processes and threads
- ✓ the O/S schedules processes and threads
- ✓ the O/S provides mechanisms for process synchronization
- ✓ the O/S provides mechanisms for interprocess communication
- ✓ the O/S provides mechanisms for deadlock handling

CPU Scheduling

- As a process executes, it changes *state*
 - **new**: The process is being created
 - **ready**: The process is waiting to be assigned to a process
 - **running**: Instructions are being executed
 - **waiting**: The process is waiting for some event to occur
 - **terminated**: The process has finished execution



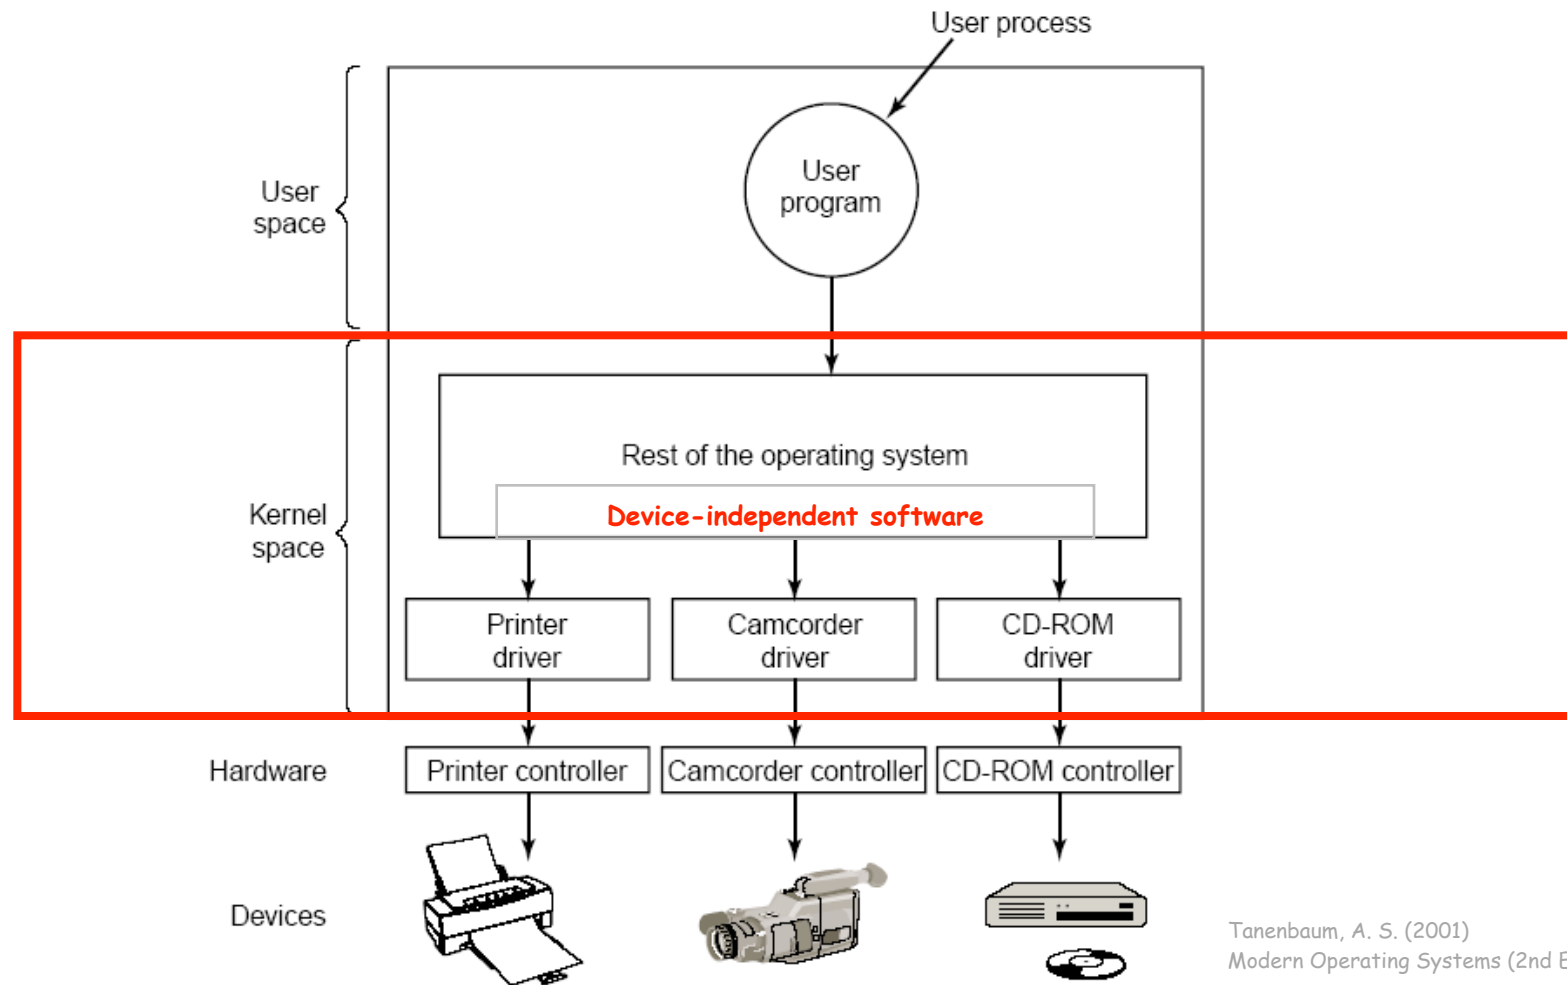
CPU Scheduling

➤ Operating System Responsibilities:

The O/S is responsible for efficiently using the CPU and providing the user with short response times

- ✓ decides which available processes in memory are to be executed by the processor
- ✓ decides what process is executed when and for how long, also reacting to external events such as I/O interrupts
- ✓ relies on a scheduling algorithm that attempts to optimize CPU utilization, throughput, latency, and/or response time, depending on the system requirements

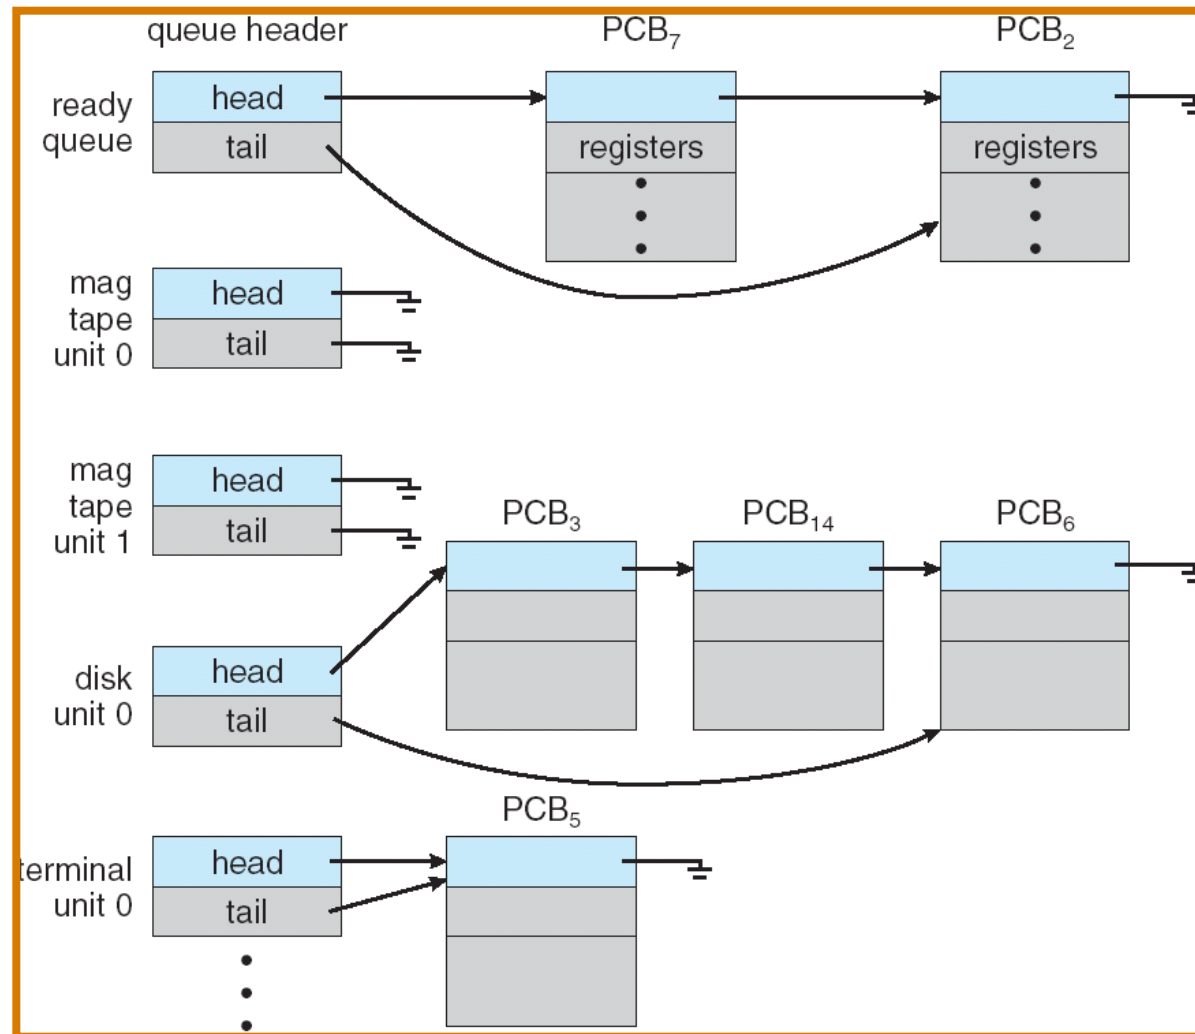
I/O Management



Tanenbaum, A. S. (2001)
Modern Operating Systems (2nd Edition).

Layers of the I/O subsystem

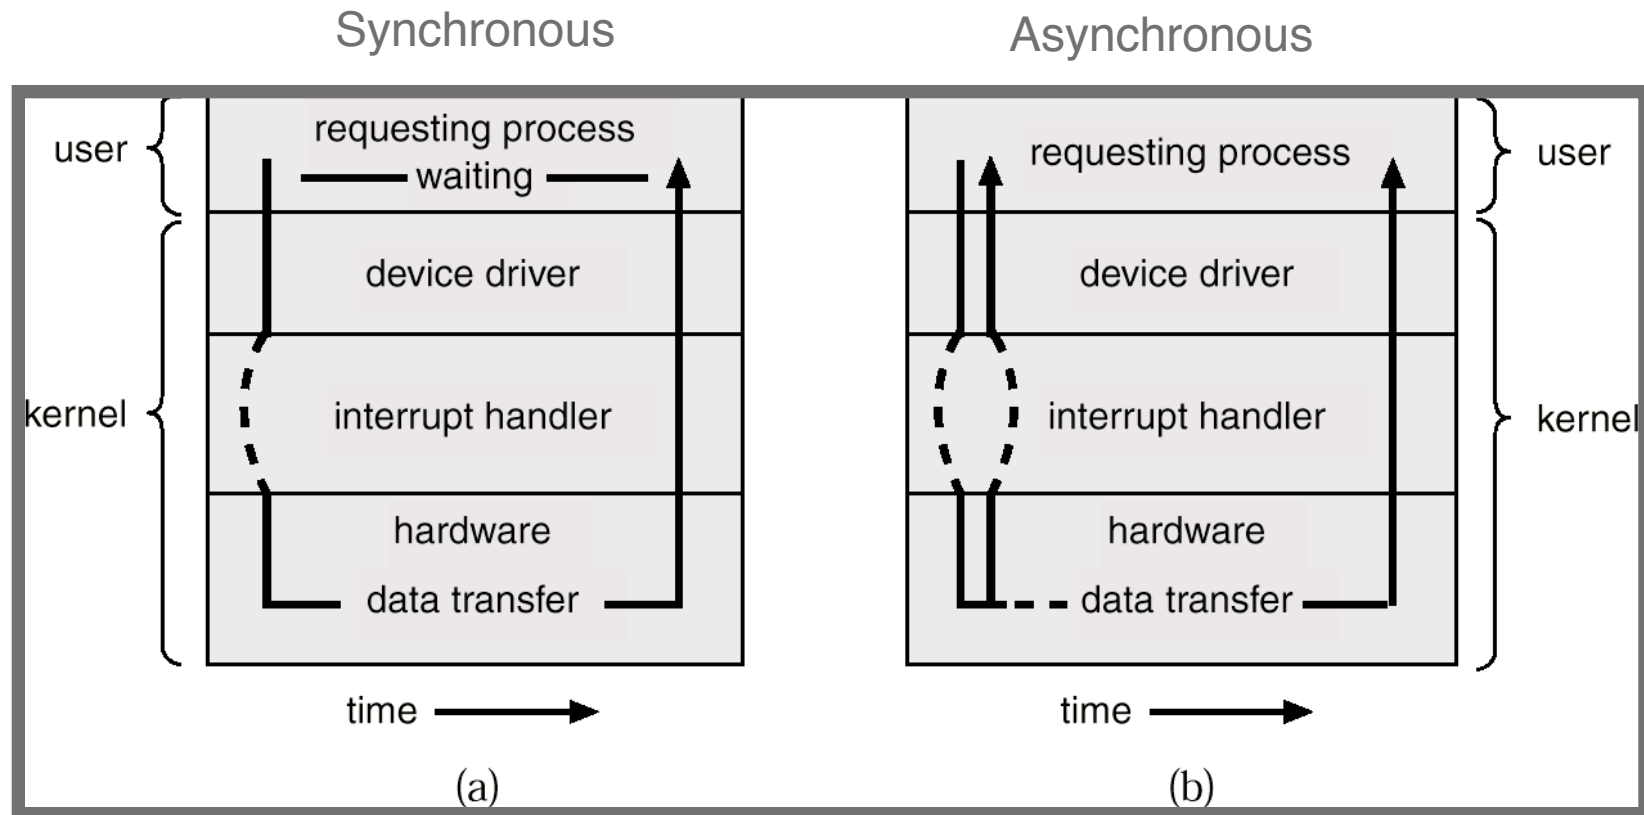
I/O Device Queues



Two I/O Methods

- After I/O starts, control returns to user program only upon I/O completion → **synchronous**
 - Wait instruction idles the CPU until the next interrupt
 - Wait loop (contention for memory access).
 - At most one I/O request is outstanding at a time, no simultaneous I/O processing.
- After I/O starts, control returns to user program without waiting for I/O completion → **asynchronous**
 - *System call* - request to the operating system to allow user to wait for I/O completion.
 - *Device-status table* contains entry for each I/O device indicating its type, address, and state.
 - Operating system indexes into I/O device table to determine device status and to modify table entry to include interrupt.

Two I/O Methods



I/O Management

➤ Operating System Responsibilities:


The O/S is responsible for controlling access to all the I/O devices

- ✓ hides the peculiarities of specific hardware devices from the user
- ✓ issues the low-level commands to the devices, catches interrupts and handles errors
- ✓ relies on software modules called "device drivers"
- ✓ provides a device-independent API to the user programs, which includes buffering

Memory Management

➤ The O/S must fit multiple processes in memory

- ✓ memory needs to be subdivided to accommodate multiple processes
- ✓ memory needs to be allocated to ensure a reasonable supply of ready processes so that the CPU is never idle
- ✓ memory management is an **optimization** task under **constraints**

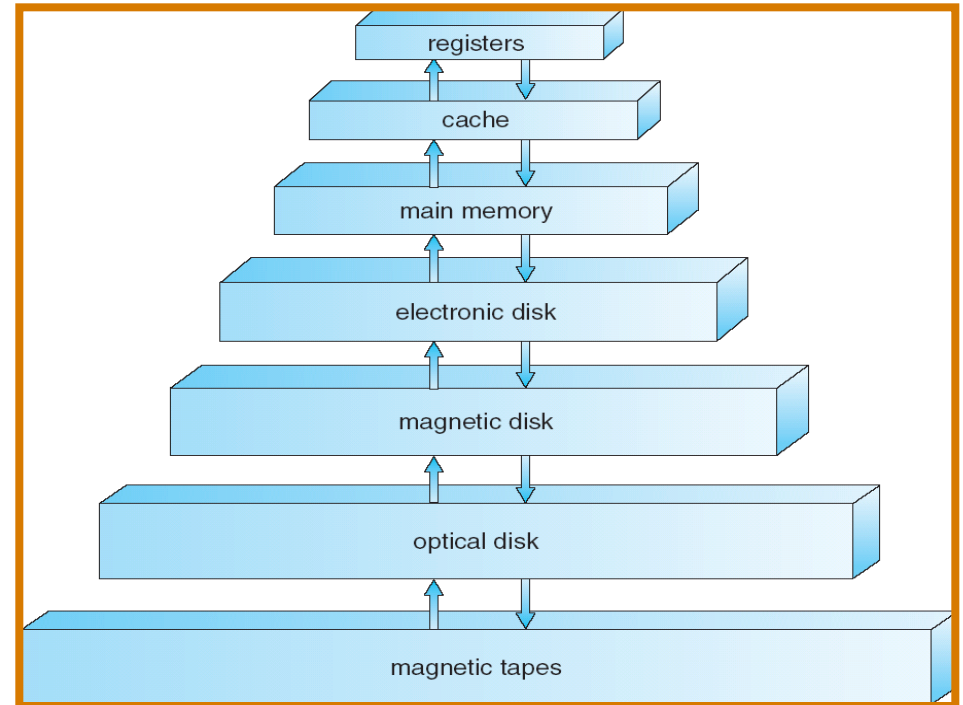


Fitting processes into memory is like fitting boxes into a fixed amount of space

Memory Management

➤ Main memory

- ✓ large array of words or bytes, each with its own address
- ✓ repository of quickly accessible data shared by the CPU and I/O devices
- ✓ volatile storage that loses its contents in case of system failure



The storage hierarchy

Performance of Various Levels of Storage

- Movement between levels of storage hierarchy can be explicit or implicit


| Level | 1 | 2 | 3 | 4 |
|---------------------------|---|-------------------------------|------------------|------------------|
| Name | registers | cache | main memory | disk storage |
| Typical size | < 1 KB | > 16 MB | > 16 GB | > 100 GB |
| Implementation technology | custom memory with multiple ports, CMOS | on-chip or off-chip CMOS SRAM | CMOS DRAM | magnetic disk |
| Access time (ns) | 0.25 – 0.5 | 0.5 – 25 | 80 – 250 | 5,000.000 |
| Bandwidth (MB/sec) | 20,000 – 100,000 | 5000 – 10,000 | 1000 – 5000 | 20 – 150 |
| Managed by | compiler | hardware | operating system | operating system |
| Backed by | cache | main memory | disk | CD or tape |

Caching

- Important principle, performed at many levels in a computer (in hardware, operating system, software)
- Information in use copied from slower to faster storage temporarily
- Faster storage (cache) checked first to determine if information is there
 - If it is, information used directly from the cache (fast)
 - If not, data copied to cache and used there
- If cache is smaller than storage being cached
 - Cache management - important design problem
 - Cache size and replacement policy

Migration of Integer A from Disk to Register

- Multitasking environments must be careful to use most recent value, not matter where it is stored in the storage hierarchy



- Multiprocessor environment must provide **cache coherency** in hardware such that all CPUs have the most recent value in their cache
- Distributed environment situation even more complex
 - Several copies of a datum can exist

Memory Management

➤ Operating System Responsibilities:

The O/S is responsible for an efficient and orderly control of storage allocation

- ✓ ensures process isolation: it keeps track of which parts of memory are currently being used and by whom
- ✓ allocates and deallocates memory space as needed: it decides which processes to load or swap out
- ✓ regulates how different processes and users can sometimes share the same portions of memory
- ✓ transfers data between main memory and disk and ensures long-term storage

OS DESIGN APPROACHES

Operating System Design and Implementation

- Start by defining goals and specifications
- Affected by choice of hardware, type of system
 - Batch, time shared, single user, multi user, distributed
- *User* goals and *System* goals
 - **User goals** - operating system should be convenient to use, easy to learn, reliable, safe, and fast
 - **System goals** - operating system should be easy to design, implement, and maintain, as well as flexible, reliable, error-free, and efficient
- No unique solution for defining the requirements of an OS
 - Large variety of solutions
 - Large variety of OS

Operating System Design and Implementation (Cont.)

- Important principle: to separate policies and mechanisms

Policy: What will be done?

Mechanism: How to do something?

- Eg. to ensure CPU protection
 - Use Timer construct (mechanism)
 - How long to set the timer (policy)
- The separation of policy from mechanism allows maximum **flexibility** if policy decisions are to be changed later

OS Design Approaches

- Simple Structure
- Layered Approach
- Microkernels
- Modules

Simple Structure

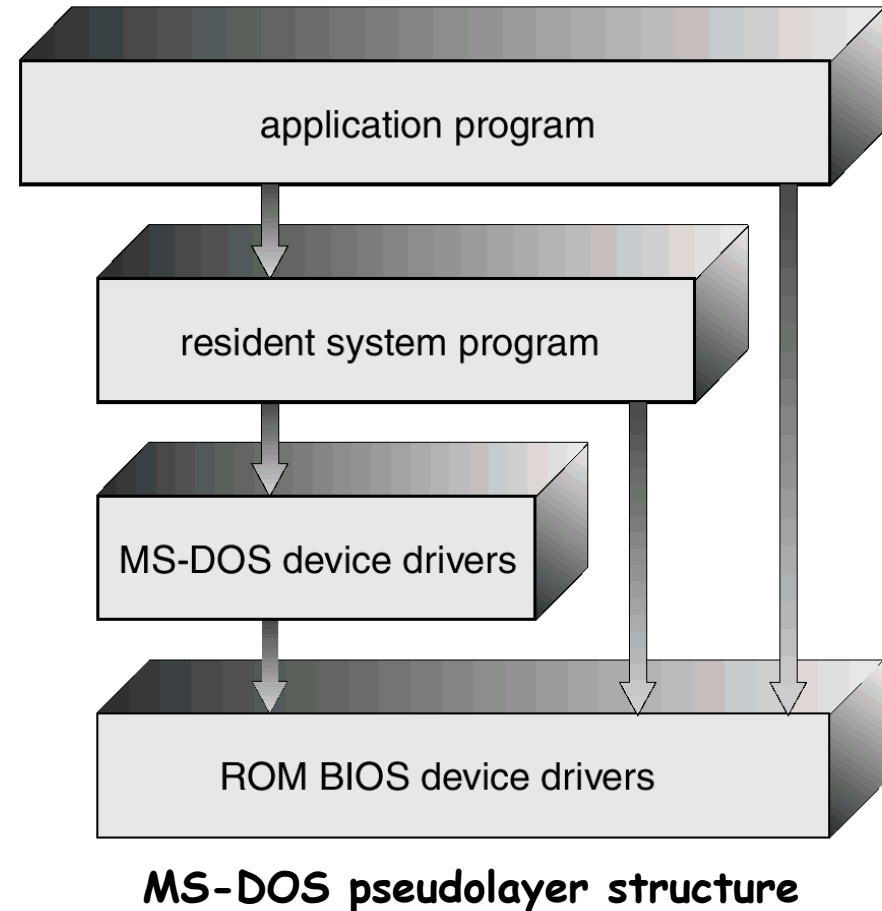
- No well defined structure
- Start as small, simple, limited systems, and then grow
- No well defined layers, not divided into modules

Simple Structure

➤ Example: MS-DOS

- ✓ initially written to provide the most functionality in the least space
- ✓ started small and grew beyond its original scope
- ✓ levels not well separated: programs could access I/O devices directly
- ✓ excuse: the hardware of that time was limited (no dual user/kernel mode)

Silberschatz, A., Galvin, P. B. and Gagne, G. (2003)
Operating Systems Concepts with Java (6th Edition).



Layered Approach

➤ Monolithic operating systems

- ✓ no one had experience in building truly large software systems
- ✓ the problems caused by mutual dependence and interaction were grossly underestimated
- ✓ such lack of structure became unsustainable as O/S grew
- ✓ Early UNIX, Linux, Windows systems --> monolithic, partially layered

➤ Enter hierarchical layers and information abstraction

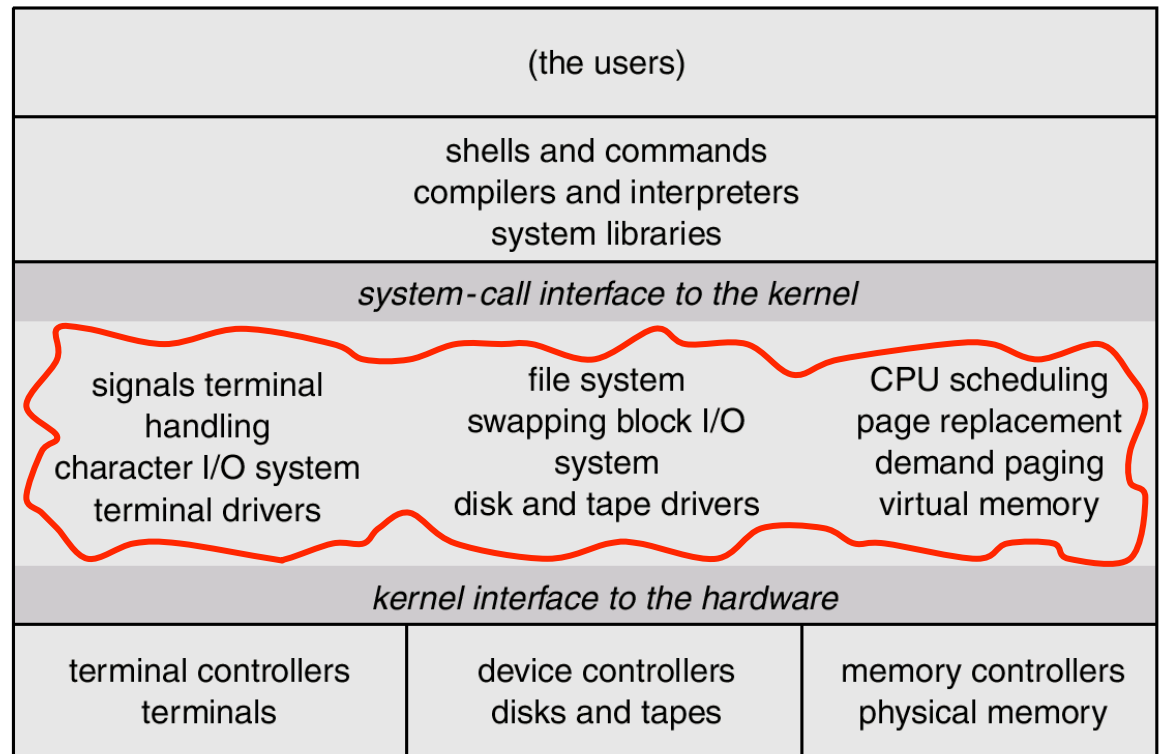
- ✓ each layer is implemented exclusively using operations provided by lower layers
- ✓ it does not need to know how they are implemented
- ✓ hence, lower layers hide the existence of certain data structures, private operations and hardware from upper layers

Simple Layered Approach

➤ The original UNIX

- ✓ enormous amount of functionality crammed into the kernel - everything below system call interface
- ✓ "The Big Mess": a collection of procedures that can call any of the other procedures whenever they need to
- ✓ no encapsulation, total visibility across the system
- ✓ very minimal layering made of thick, monolithic layers


Silberschatz, A., Galvin, P. B. and Gagne, G. (2003)
Operating Systems Concepts with Java (6th Edition).



UNIX system structure

Full Layered Approach


- The operating system is divided into a number of layers (levels), each built on top of lower layers.
 - The bottom layer (layer 0), is the hardware;
 - The highest (layer N) is the user interface.
- With modularity, layers are selected such that each uses functions (operations) and services of only lower-level layers
- THE system (by Dijkstra), MULTICS, GLUnix, VAX/VMS



Layered Approach

- Layers can be debugged and replaced independently without bothering the other layers above and below

- ✓ famous example of strictly layered architecture: the TCP/IP networking stack



Layered Approach


Theoretical model of operating system design hierarchy

| | Level | Name | Objects | Example Operations |
|----------|-------|-----------------------|---|---|
| shell | 13 | Shell | User programming environment | Statements in shell language |
| | 12 | User processes | User processes | Quit, kill, suspend, resume |
| O/S | 11 | Directories | Directories | Create, destroy, attach, detach, search, list |
| | 10 | Devices | External devices, such as printers, displays, and keyboards | Open, close, read, write |
| | 9 | File system | Files | Create, destroy, open, close, read, write |
| | 8 | Communications | Pipes | Create, destroy, open, close, read, write |
| | 7 | Virtual memory | Segments, pages | Read, write, fetch |
| | 6 | Local secondary store | Blocks of data, device channels | Read, write, allocate, free |
| | 5 | Primitive processes | Primitive processes, semaphores, ready list | Suspend, resume, wait, signal |
| hardware | 4 | Interrupts | Interrupt-handling programs | Invoke, mask, unmask, retry |
| | 3 | Procedures | Procedures, call stack, display | Mark stack, call, return |
| | 2 | Instruction set | Evaluation stack, microprogram interpreter, scalar and array data | Load, store, add, subtract, branch |
| | 1 | Electronic circuits | Registers, gates, buses, etc. | Clear, transfer, activate, complement |

Stallings, W. (2004) *Operating Systems: Internals and Design Principles* (5th Edition).

Layered Approach

- Major difficulty with layering
 - ✓ . . . appropriately defining the various layers!
 - ✓ layering is only possible if all function dependencies can be sorted out into a Directed Acyclic Graph (DAG)
 - ✓ however there might be conflicts in the form of circular dependencies (“cycles”)



Circular dependency on top of a DAG

Layered Approach

➤ Circular dependencies in an O/S organization

- ✓ example: disk driver routines vs. CPU scheduler routines
 - the device driver for the backing store (disk space used by virtual memory) may need to wait for I/O, thus invoke the CPU-scheduling layer
 - the CPU scheduler may need the backing store driver for swapping in and out parts of the table of active processes

➤ Other difficulty: efficiency


- ✓ the more layers, the more indirections from function to function and the bigger the overhead in function calls
- ✓ backlash against strict layering: return to fewer layers with more functionality

Microkernel System Structure


➤ The microkernel approach

- ✓ a microkernel is a reduced operating system core that contains only essential O/S functions
- ✓ the idea is to minimize the kernel by moving up as much functionality as possible from the kernel into user space
- ✓ many services traditionally included in the O/S are now external subsystems running as user processes
 - device drivers
 - file systems
 - virtual memory manager
 - windowing system
 - security services, etc.

Layered OS vs Microkernel



(a) Layered kernel



(b) Microkernel

Microkernel System Structure

➤ Benefits of the microkernel approach

- ✓ **extensibility** — it is easier to extend a microkernel-based O/S as new services are added in user space, not in the kernel
- ✓ **portability** — it is easier to port to a new CPU, as changes are needed only in the microkernel, not in the other services
- ✓ **reliability & security** — much less code is running in kernel mode; failures in user-space services don't affect kernel space

➤ Detriments of the microkernel approach

- ✓ again, performance overhead due to communication from user space to kernel space
- ✓ not always realistic: some functions (I/O) must remain in kernel space, forcing a separation between “policy” and “mechanism”

- Examples: QNX, Tru64 UNIX, Mach (CMU), Windows NT

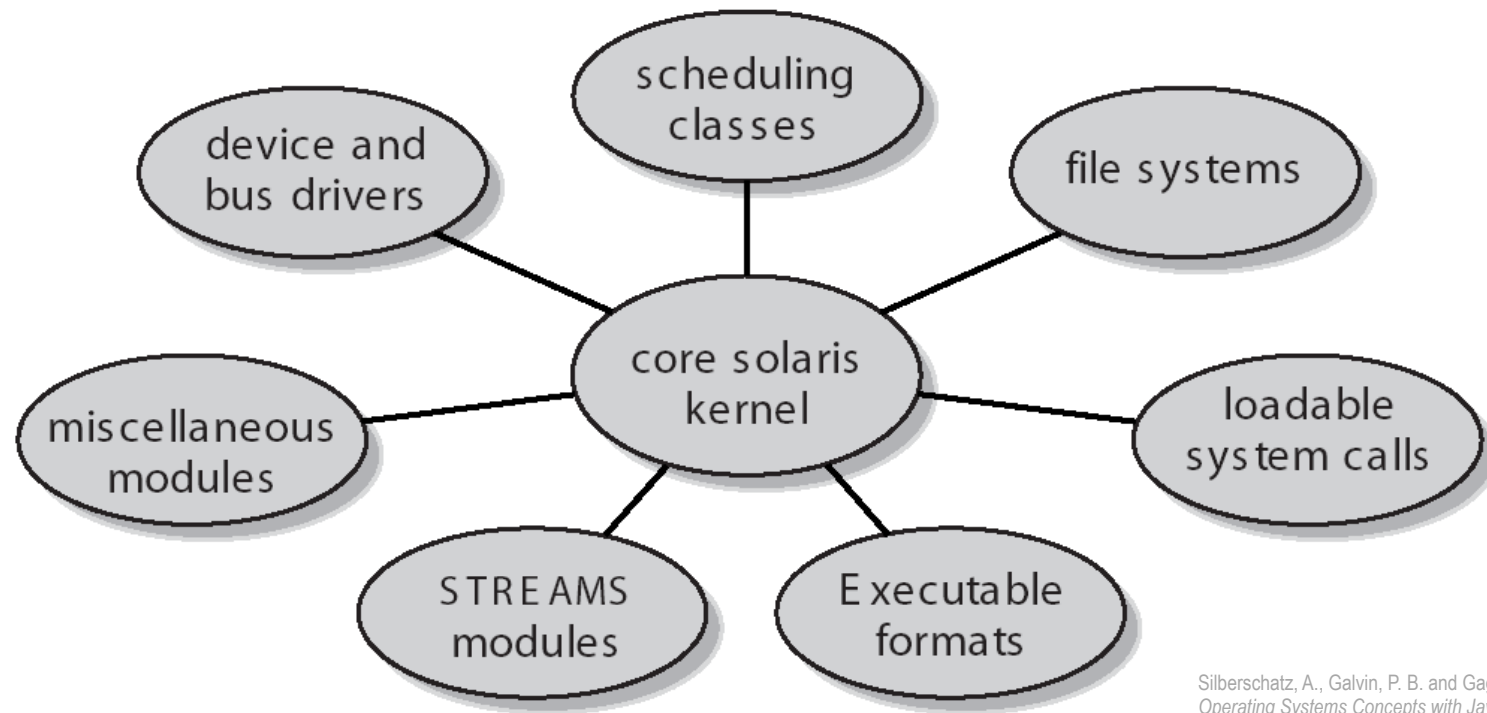
Modular Approach

➤ The modular approach

- ✓ many modern operating systems implement kernel **modules** (Modern UNIX, Solaris, Linux, Windows, Mac OS X)
- ✓ this is similar to the object-oriented approach:
 - each core component is separate
 - each talks to the others over known interfaces
 - each is loadable as needed within the kernel
- ✓ overall, modules are similar to layers but with more flexibility (any model could call any other module)
- ✓ modules are also similar to the microkernel approach, except they are inside the kernel and don't need message passing

Modular Approach

- Modules are used in Solaris, Linux and Mac OS X



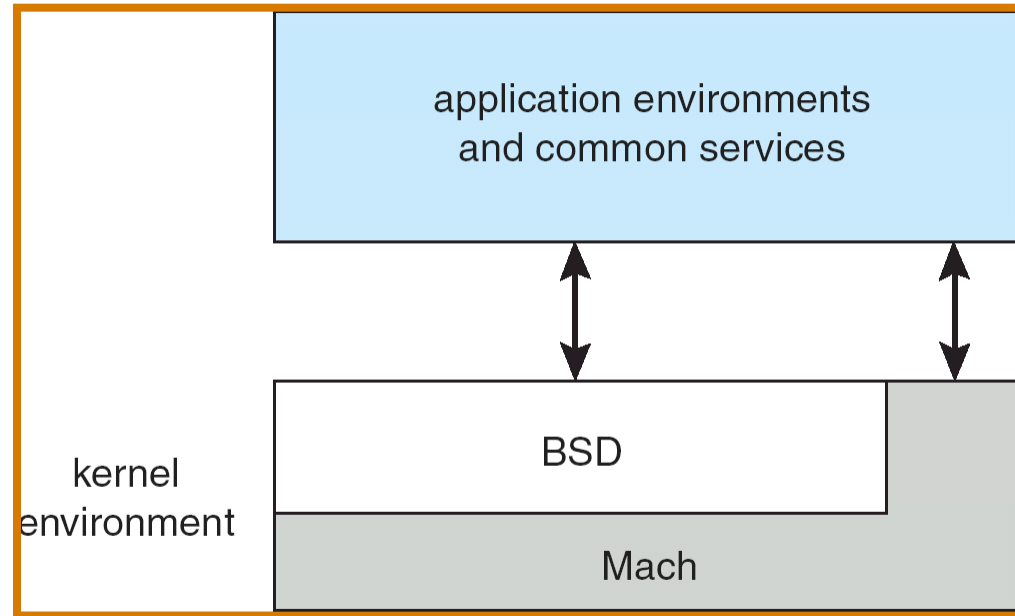
Silberschatz, A., Galvin, P. B. and Gagne, G. (2003)
Operating Systems Concepts with Java (6th Edition).

The Solaris loadable modules

Hybrid Systems

- Many real OS use combination of different approaches
- **Linux:** monolithic & modular
- **Windows:** monolithic & microkernel & modular
- **Mac OS X:** microkernel & modular


Mac OS X Structure - Hybrid



- **BSD:** provides support for command line interface, networking, file system, POSIX API and threads
- **Mach:** memory management, RPC, IPC, message passing

Summary

- OS Design Approaches
 - Mechanism vs Policy
 - Monolithic Systems,
 - Layered Approach, Microkernels, Modules
- Major OS Components
 - Processes
 - CPU Scheduling
 - I/O Management
 - Memory management



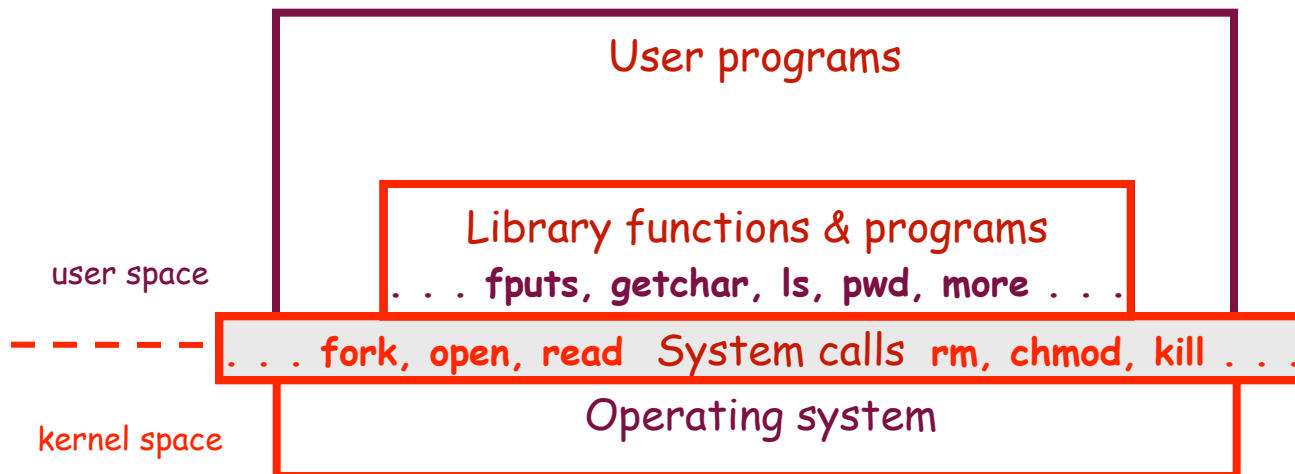
- Reading Assignment: Chapter 3 from Silberschatz.

Acknowledgements

- “Operating Systems Concepts” book and supplementary material by A. Silberschatz, P. Galvin and G. Gagne
- “Operating Systems: Internals and Design Principles” book and supplementary material by W. Stallings
- “Modern Operating Systems” book and supplementary material by A. Tanenbaum
- R. Doursat and M. Yuksel from UNR

System Calls

- System calls are the only entry points into the kernel and system

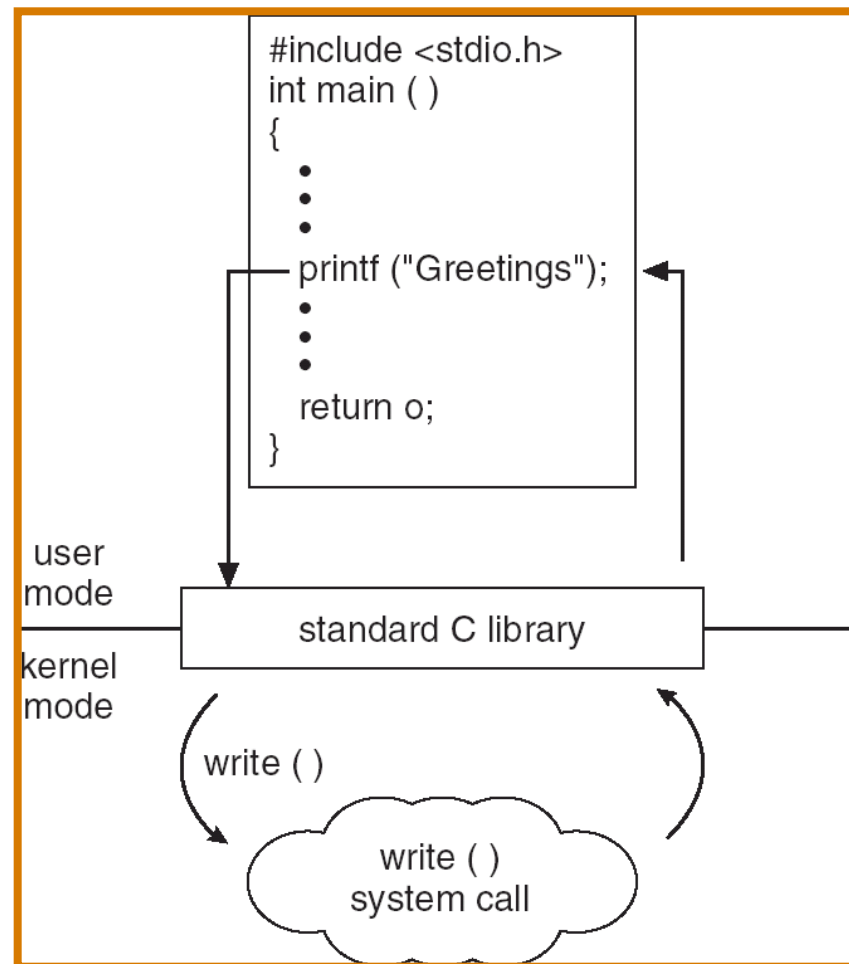


the "middleman's counter"

- Programming interface to the services provided by the OS
- All programs needing resources must use system calls
- Most UNIX commands are actually library functions and utility programs (e.g., shell interpreter) built on top of the system calls

Example

- C program invoking printf() library call, which calls write() system call

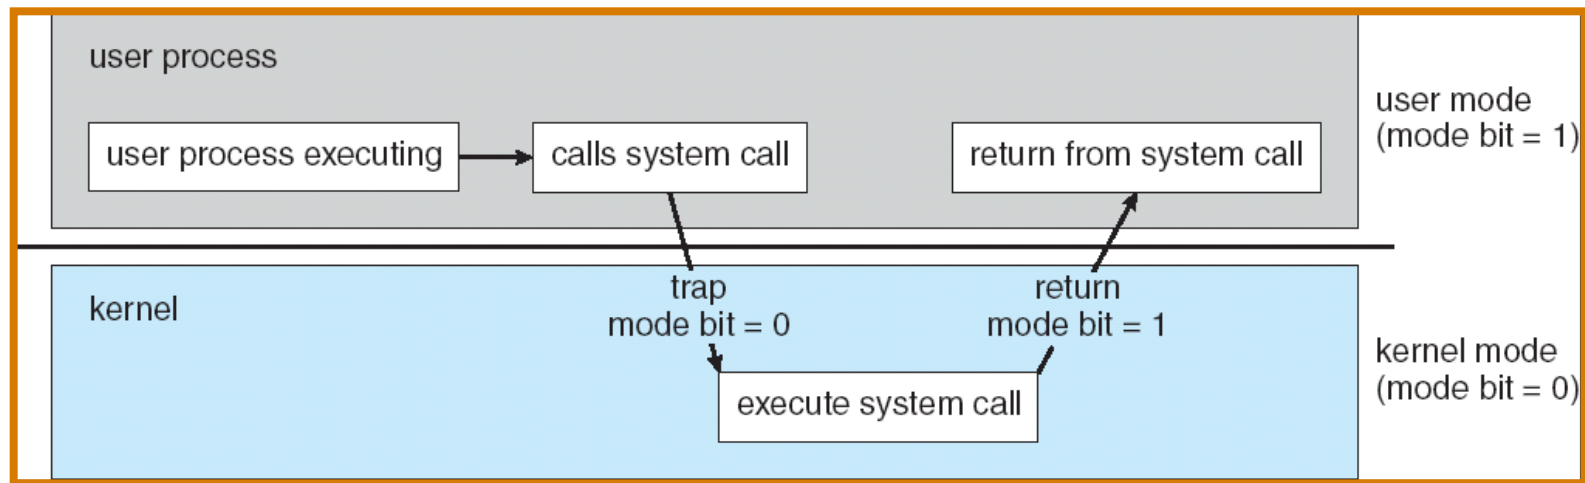


Dual-Mode Operation

- **Dual-mode** operation allows OS to protect itself and other system components
 - **User mode** and **kernel mode**
 - **Mode bit** provided by hardware
 - Provides ability to distinguish when system is running user code or kernel code
 - Protects OS from errant users, and errant users from each other
 - Some instructions designated as **privileged**, only executable in kernel mode
 - System call changes mode to kernel, return from call resets it to user

Transition from User to Kernel Mode

- How to prevent user program getting stuck in an infinite loop / process hogging resources
 - ➔ **Timer:** Set interrupt after specific period (1ms to 1sec)
 - Operating system decrements counter
 - When counter zero generate an interrupt
 - Set up before scheduling process to regain control or terminate program that exceeds allotted time



Questions

- At the system boot time, what should be the mode of operation?
- When to switch to user mode?
- When to switch to kernel mode?
- Which of these are mechanisms?
- Which of these are policies?