

CSE 421/521 - Operating Systems
Fall 2014

LECTURE - V
CPU SCHEDULING - I

Tevfik Koşar

University at Buffalo
September 9th, 2014

Roadmap

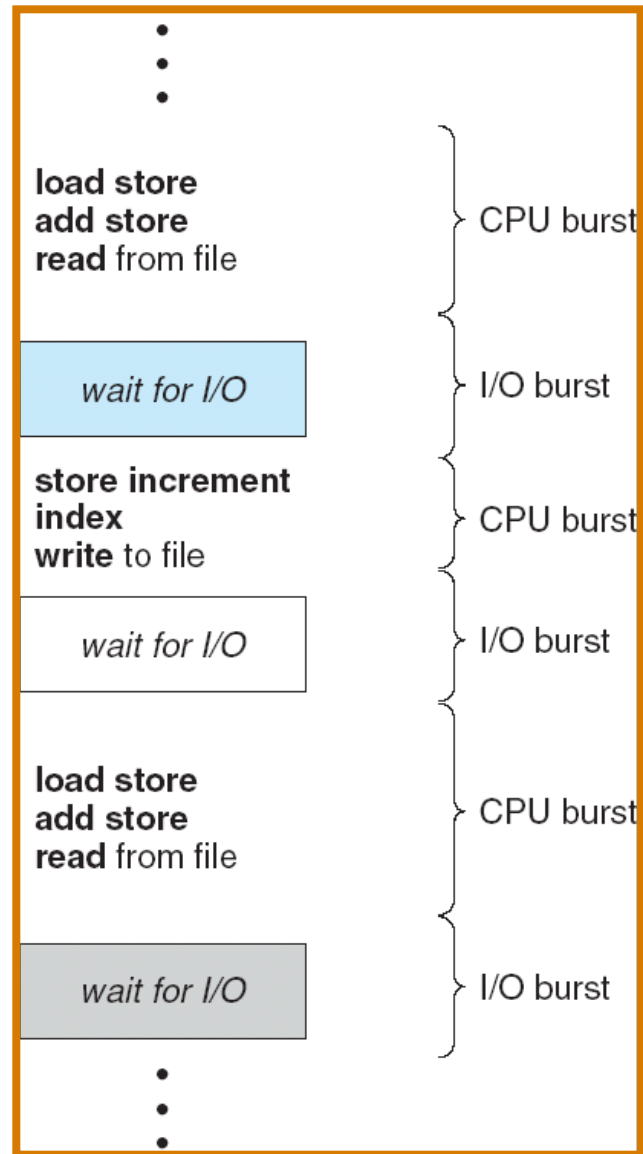
- CPU Scheduling
 - Basic Concepts
 - Scheduling Criteria & Metrics
 - Different Scheduling Algorithms
 - FCFS
 - SJF
 - Priority
 - RR
 - Preemptive vs Non-preemptive Scheduling
 - Gantt Charts & Performance Comparison
 - Multilevel Feedback Queues




Basic Concepts

- Multiprogramming is needed for efficient CPU utilization
- **CPU Scheduling:** deciding which processes to execute when
- Process execution begins with a **CPU burst**, followed by an **I/O burst**
- CPU-I/O Burst Cycle - Process execution consists of a *cycle* of CPU execution and I/O wait

Alternating Sequence of CPU And I/O Bursts




Histogram of CPU-burst Durations



Process State

- As a process executes, it changes *state*
 - **new**: The process is being created
 - **ready**: The process is waiting to be assigned to a process
 - **running**: Instructions are being executed
 - **waiting**: The process is waiting for some event to occur
 - **terminated**: The process has finished execution



CPU Scheduler

- Selects from among the processes in memory that are ready to execute, and allocates the CPU to one of them
 - ➔ short-term scheduler
- CPU scheduling decisions may take place when a process:
 1. Switches from running to waiting state
 2. Switches from running to ready state
 3. Switches from waiting to ready
 4. Terminates
 5. A new process arrives
- Scheduling under 1 and 4 is *nonpreemptive/cooperative*
 - Once a process gets the CPU, keeps it until termination/switching to waiting state/release of the CPU
- All other scheduling is *preemptive*
 - Most OS use this
 - Cost associated with access to shared data
 - i.e. time quota expires

Dispatcher

- Dispatcher module gives control of the CPU to the process selected by the short-term scheduler;
Its function involves:
 - switching context
 - switching to user mode
 - jumping to the proper location in the user program to restart that program
- *Dispatch latency* - time it takes for the dispatcher to stop one process and start another running

Scheduling Criteria

- **CPU utilization** - keep the CPU as busy as possible
--> **maximize**
- **Throughput** - # of processes that complete their execution per time unit --> **maximize**
- **Turnaround time** - amount of time passed to finish execution of a particular process --> **minimize**
 - i.e. execution time + waiting time
- **Waiting time** - total amount of time a process has been waiting in the ready queue --> **minimize**
- **Response time** - amount of time it takes from when a request was submitted until the first response is produced, **not** output (for time-sharing environment) --> **minimize**

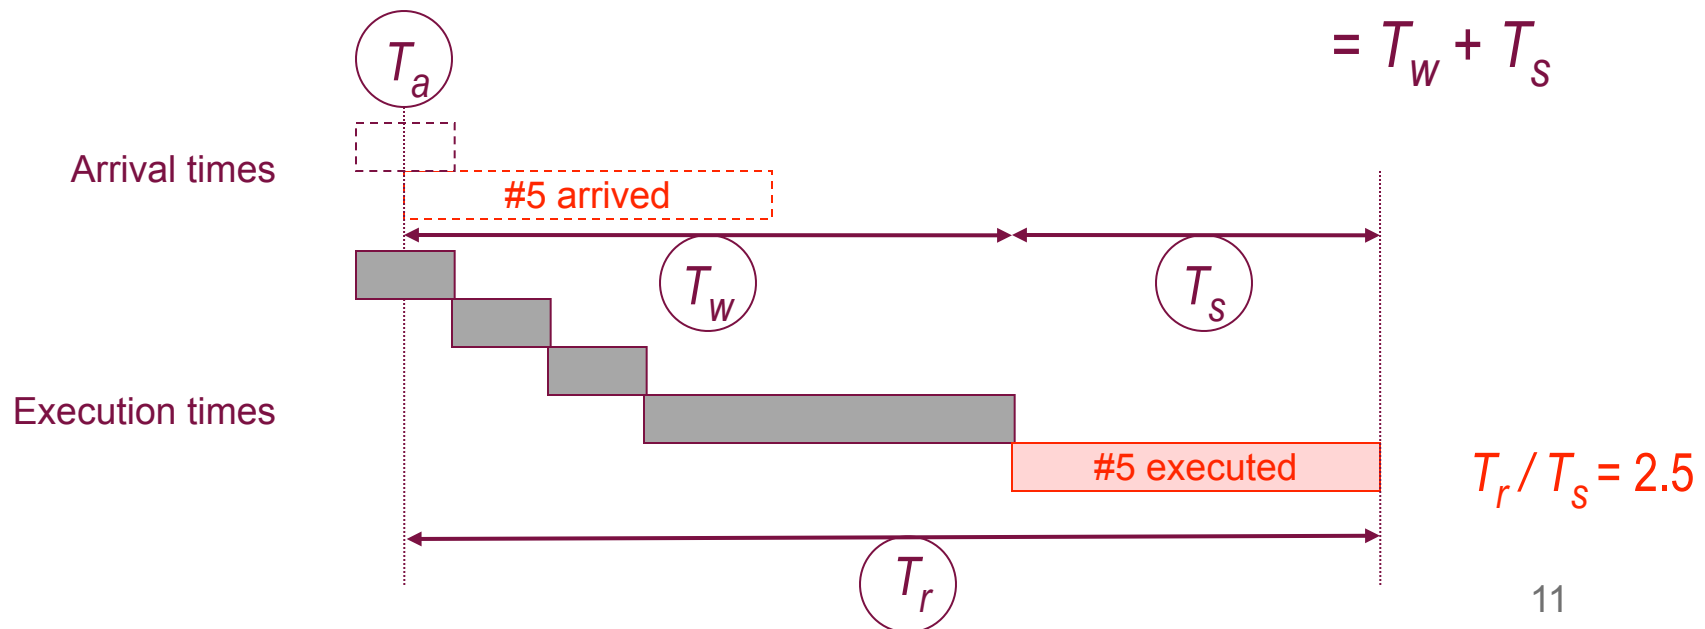
Optimization Criteria

- Maximize CPU utilization
- Maximize throughput
- Minimize turnaround time
- Minimize waiting time
- Minimize response time

Scheduling Metrics


➤ Scheduling metrics

- ✓ arrival time T_a = time the process became “Ready” (again)
- ✓ wait time T_w = time spent waiting for the CPU
- ✓ service time T_s = time spent executing in the CPU
- ✓ turnaround time T_r = total time spent waiting and executing



First-Come, First-Served (FCFS) Scheduling

- ✓ processes are assigned the CPU in the order they request it
- ✓ when the running process blocks, the first “Ready” is run next
- ✓ when a process gets “Ready”, it is put at the end of the queue



FCFS	Finish Time	A	3	B	9	C	13	D	18	E	20	Mean
	Turnaround Time (T_r)		3		7		9		12		12	8.60
	T_r/T_s		1.00		1.17		2.25		2.40		6.00	2.56

FCFS scheduling policy


Stallings, W. (2004) *Operating Systems: Internals and Design Principles* (5th Edition).

FCFS Scheduling - Example

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

- Suppose that the processes arrive in the order: P_1 , P_2 , P_3

The **Gantt Chart** for the schedule is:




- Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- Average waiting time: $(0 + 24 + 27)/3 = 17$

FCFS Scheduling - Example

Suppose that the processes arrive in the order

$$P_2, P_3, P_1$$

- The Gantt chart for the schedule is:




- Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- Average waiting time: $(6 + 0 + 3)/3 = 3$
- Much better than previous case
- *Convoy effect* short process behind long process

Shortest-Job-First (SJF) Scheduling

- Associate with each process the length of its next CPU burst. Use these lengths to schedule the process with the shortest time
- Two schemes:
 - **nonpreemptive** - once CPU given to the process it cannot be preempted until completes its CPU burst
 - **preemptive** - if a new process arrives with CPU burst length less than remaining time of current executing process, preempt.
-->This scheme is know as the **Shortest-Remaining-Time-First (SRTF)**
- SJF is optimal - gives minimum average waiting time for a given set of processes

Non-Preemptive SJF

- ✓ nonpreemptive, assumes the run times are known in advance
- ✓ among several equally important “Ready” jobs (or CPU bursts), the scheduler picks the one that will finish the earliest




SJF scheduling policy

Stallings, W. (2004) *Operating Systems: Internals and Design Principles (5th Edition)*.

Non-Preemptive SJF - Example

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

- SJF (non-preemptive) **Gantt Chart**




- Average waiting time = $(0 + 6 + 3 + 7)/4 = 4$

Preemptive SJF (SRT)

➤ Shortest Remaining Time (SRT)

- ✓ preemptive version of SJF, also assumes known run time
- ✓ choose the process whose remaining run time is shortest
- ✓ allows new short jobs to get good service



SRT	Finish Time	A	3	B	15	C	8	D	20	E	10	Mean
	Turnaround Time (T_r)		3		13		4		14		2	7.20
	T_r/T_s		1.00		2.17		1.00		2.80		1.00	1.59


SRT scheduling policy

Stallings, W. (2004) *Operating Systems: Internals and Design Principles* (5th Edition).

Example of Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

- SJF (preemptive) **Gantt Chart**



Priority Scheduling

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer \equiv highest priority)
 - Preemptive
 - nonpreemptive
- SJF is a priority scheduling where priority is the predicted next CPU burst time
- Problem \equiv **Starvation** - low priority processes may never execute
- Solution \equiv **Aging** - as time progresses increase the priority of the process

Example of Priority

	<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>	<u>Priority</u>
-	P_1	0.0	7	2
	P_2	2.0	4	1
	P_3	4.0	1	4
	P_4	5.0	4	3


- Priority (non-preemptive)
 - $P_1 \rightarrow P_2 \rightarrow P_4 \rightarrow P_3$
- Priority (preemptive)
 - ??

Round Robin (RR)

- Each process gets a small unit of CPU time (*time quantum*), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units at once. No process waits more than $(n-1)q$ time units.
- Performance
 - q large \Rightarrow FIFO
 - q small $\Rightarrow q$ must be large with respect to context switch, otherwise overhead is too high

Round Robin (RR)

- ✓ preemptive FCFS, based on a timeout interval, the **quantum** q
- ✓ the running process is interrupted by the clock and put last in a FIFO “Ready” queue; then, the first “Ready” process is run instead




RR $q = 1$	Finish Time	A	4	B	18	C	17	D	20	E	15	Mean
	Turnaround Time (T_r)		4		16		13		14		7	10.80
	T_r/T_s		1.33		2.67		3.25		2.80		3.50	2.71

RR ($q = 1$) scheduling policy

Stallings, W. (2004) *Operating Systems: Internals and Design Principles (5th Edition)*.

Round Robin (RR)

- ✓ a crucial parameter is the quantum q (generally $\sim 10\text{--}100\text{ms}$)
 - q should be big compared to context switch latency ($\sim 10\mu\text{s}$)
 - q should be less than the longest CPU bursts, otherwise RR degenerates to FCFS




RR ($q = 4$) scheduling policy

Stallings, W. (2004) *Operating Systems: Internals and Design Principles* (5th Edition).

Example of RR with Time Quantum = 20


<u>Process</u>	<u>Burst Time</u>
P_1	53
P_2	17
P_3	68
P_4	24

- For $q=20$, the **Gantt chart** is:




Typically, higher average turnaround than SJF,
but better *response*

Time Quantum and Context Switch Time




Turnaround Time Varies With The Time Quantum



MULTILEVEL QUEUES

Multilevel Queue Scheduling




Multilevel Feedback Queue

- A process can move between the various queues; aging can be implemented this way
- Multilevel-feedback-queue scheduler defined by the following parameters:
 - number of queues
 - scheduling algorithms for each queue
 - method used to determine when to upgrade a process
 - method used to determine when to demote a process
 - method used to determine which queue a process will enter when that process needs service

Example of Multilevel Feedback Queue

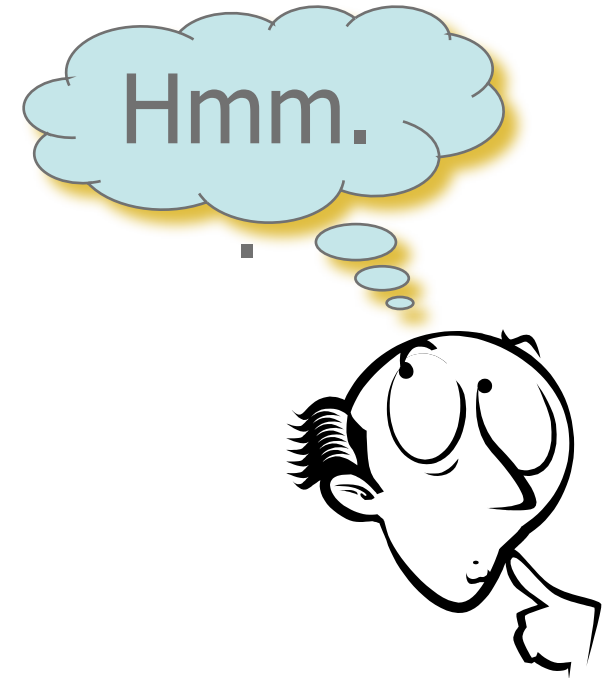
- Three queues:
 - Q_0 - RR with time quantum 8 milliseconds
 - Q_1 - RR time quantum 16 milliseconds
 - Q_2 - FCFS
- Scheduling
 - A new job enters queue Q_0 which is served FCFS. When it gains CPU, job receives 8 milliseconds. If it does not finish in 8 milliseconds, job is moved to queue Q_1 .
 - At Q_1 job is again served FCFS and receives 16 additional milliseconds. If it still does not complete, it is preempted and moved to queue Q_2 .

Multilevel Feedback Queues



Summary

- CPU Scheduling
 - Basic Concepts
 - Scheduling Criteria & Metrics
 - Different Scheduling Algorithms
 - FCFS
 - SJF
 - Priority
 - RR
 - Multilevel Feedback Queues



- Next Lecture: Project-1 Discussion
- Reading Assignment: Chapter 5 from Silberschatz.

Acknowledgements

- “Operating Systems Concepts” book and supplementary material by A. Silberschatz, P. Galvin and G. Gagne
- “Operating Systems: Internals and Design Principles” book and supplementary material by W. Stallings
- “Modern Operating Systems” book and supplementary material by A. Tanenbaum
- R. Doursat and M. Yuksel from UNR