

CSE 421/521 - Operating Systems  
Fall 2014

LECTURE - VI

# PROJECT-1 DISCUSSION

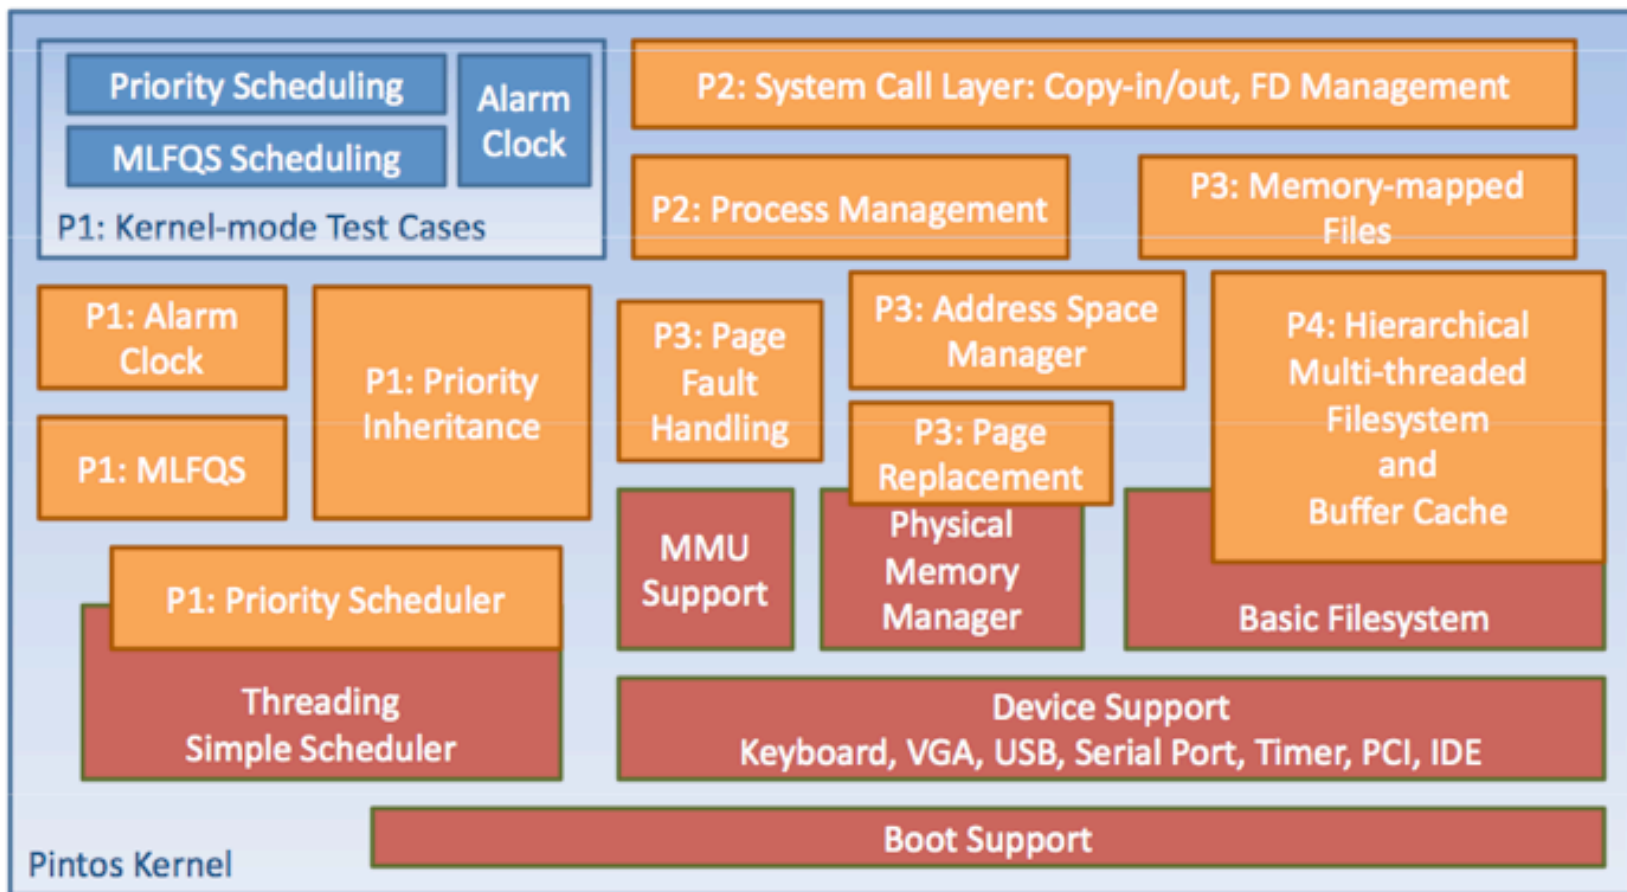

Tevfik Koşar

University at Buffalo  
September 11<sup>th</sup>, 2014

# Pintos Projects

1. Threads <-- CSE 421/521 Project 1
2. User Programs
3. Virtual Memory
4. File Systems

# Pintos after full implementation (post prj-4)



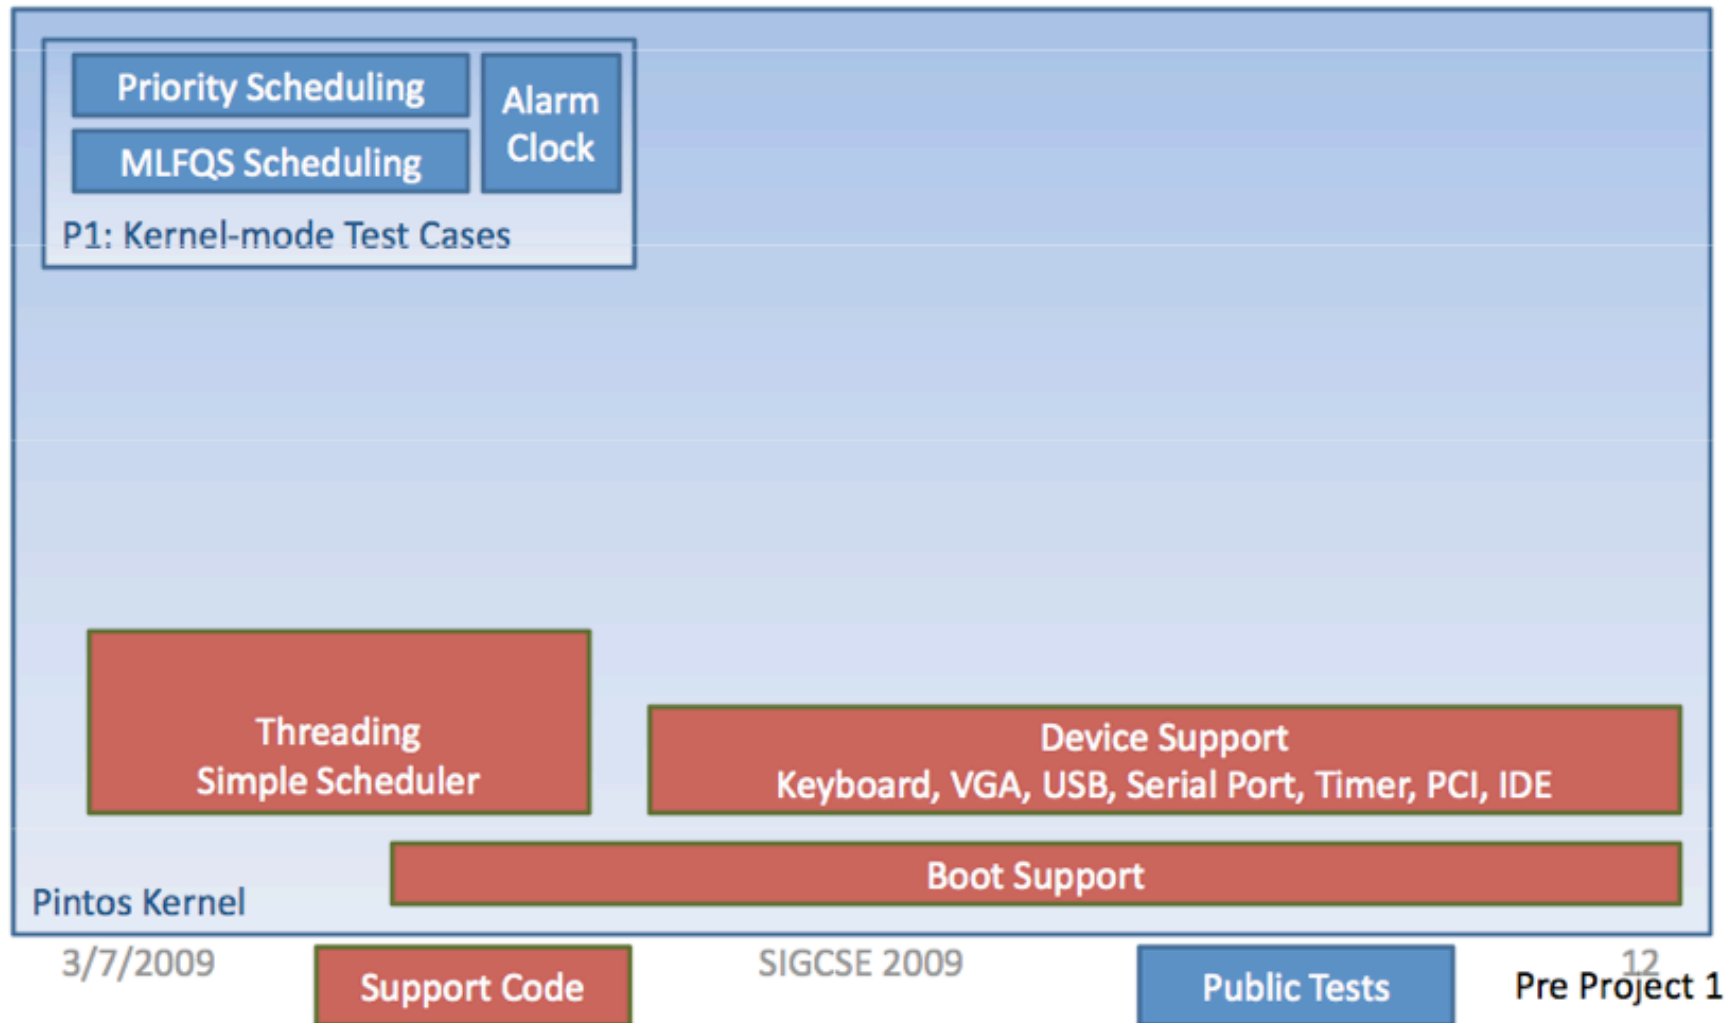
Support Code

Students Create

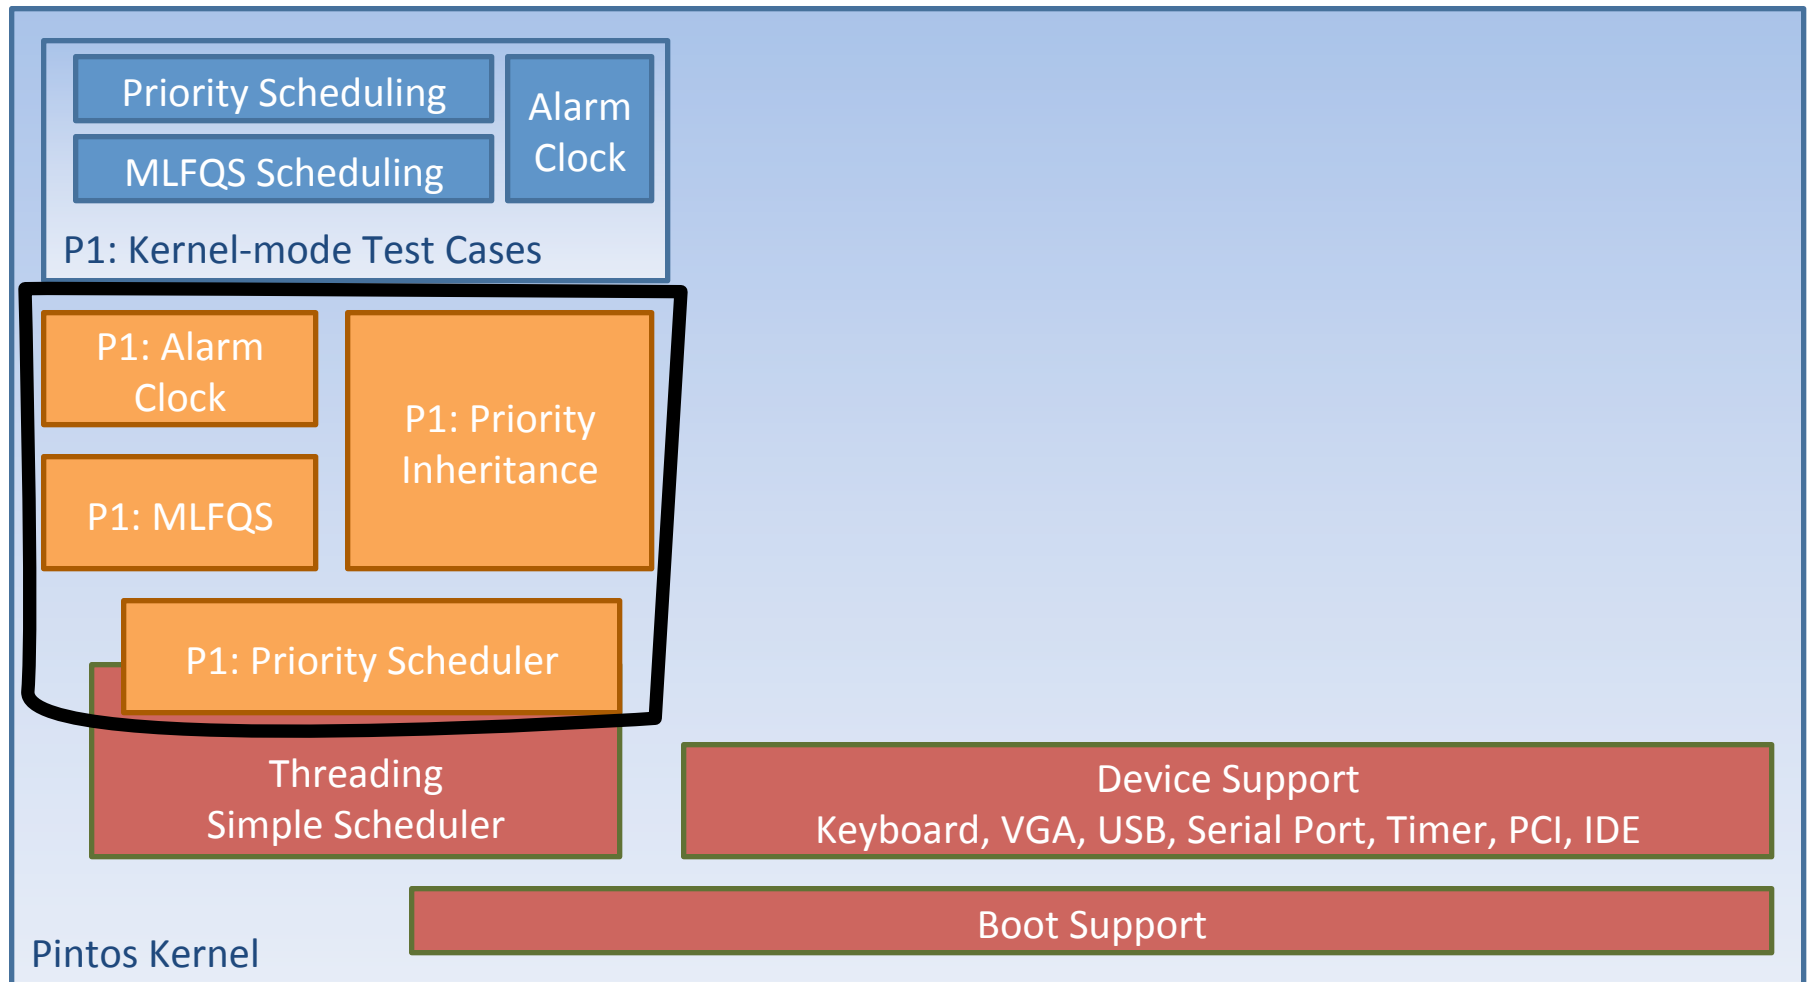
Public Tests

Post Project 4

# Yo will be provided with this (pre prj-1)



# You will implement this (post prj-1)



# Step 1: Preparation

## READ:

- Chapters 3-5 from Silberschatz
- Lecture slides on Processes, Threads and CPU Scheduling
- From Pintos Documentation:
  - Section 1 - Introduction
  - Section 2 - Threads
  - Appendix A1- Pintos Loading
  - Appendix A2 - Threads
  - Appendix A3 - Synchronization
  - Appendix B - 4.4BSD Scheduler

## Step 2: Setting Up Pintos

- Option 1: Fetch source code from:  
[/web/faculty/tkosar/cse421-521/projects/project-1/pintos.tar](http://web/faculty/tkosar/cse421-521/projects/project-1/pintos.tar)  
then follow the setup guidelines we have provided on Piazza  
==> tested on dragonforce, styx, nickelback
- Option 2: Use the Pintos VM we have prepared for you:  
<http://ftp.cse.buffalo.edu/CSE421/Pintos.ova>  
requires Virtualbox software  
==> will work on most Linux, Windows, Mac systems
- Option 3: Grab the source directly from pintos git repository:  
<http://pintos-os.org/cgi-bin/gitweb.cgi>  
you would also need Bochs or Qemu simulator with this option

## Step 3: Implementation

1. Alarm Clock
2. Priority Scheduler (with priority donation)
3. Multilevel Feedback Queue Scheduler




# Pintos Thread System

Defined in threads/thread.h:

```
struct thread
{
    tid_t tid;          /* Thread identifier. */
    enum thread_status status; /* Thread state. */
    char name[16]; /* Name (for debugging purposes). */
    uint8_t *stack; /* Saved stack pointer. */
    int priority; /* Priority. */
    struct list_elem allelem; /* List element for all-threads list. */
    /* Shared between thread.c and synch.c. */
    struct list_elem elem; /* List element. */
};
```

You add more fields here as you need them.

```
#ifdef USERPROG
    /* Owned by userprog/process.c. */
    uint32_t *pagedir; /* Page directory. */
#endif
    /* Owned by thread.c. */
    unsigned magic; /* Detects stack overflow. */
};
```



# Pintos Thread System

- Read threads/thread.c and threads/synch.c to understand
  - How the switching between threads occur
  - How the provided scheduler works
  - How the various synchronizations primitives work

# Task 1: Alarm Clock

- Reimplement `timer_sleep()` in `devices/timer.c` without busy waiting

`/* Suspends execution for approximately TICKS timer ticks. */`

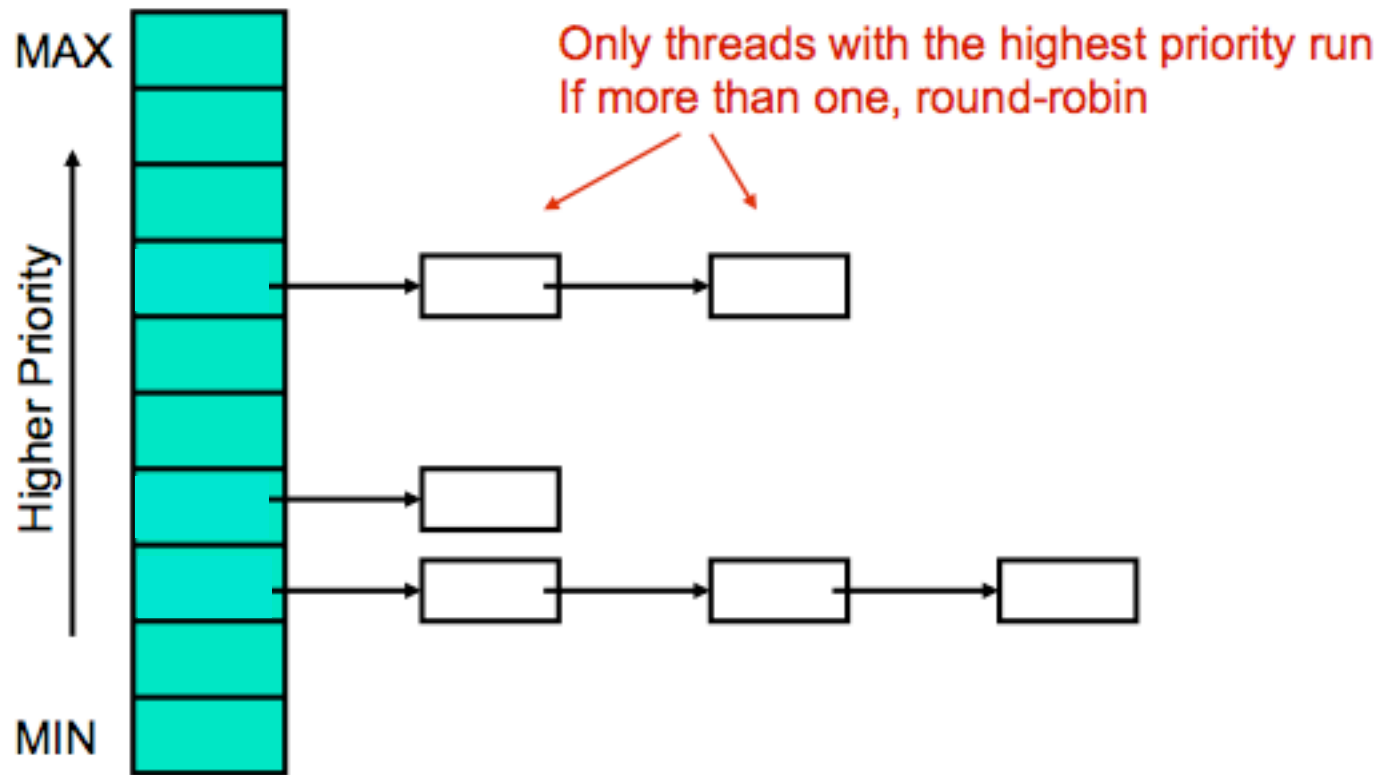
```
void timer_sleep (int64_t ticks){  
    int64_t start = timer_ticks ();  
    ASSERT (intr_get_level () == INTR_ON);  
    while (timer_elapsed (start) < ticks)  
        thread_yield ();  
}
```

- Implementation details
  - Remove thread from ready list and put it back after sufficient ticks have elapsed

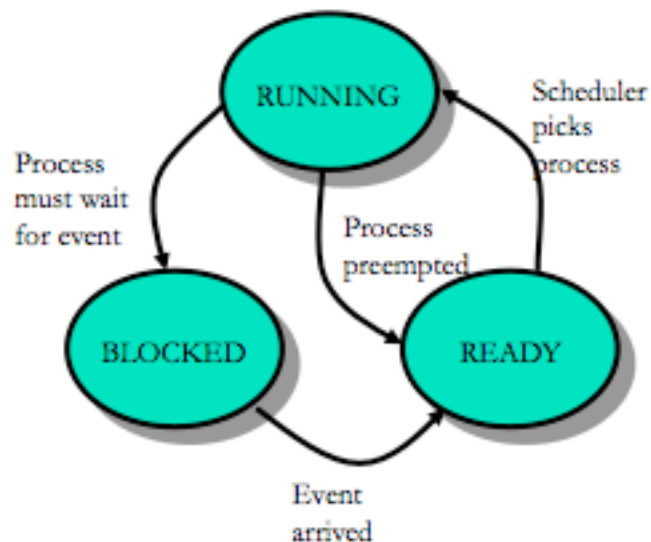
## Task 2A: Priority Scheduler

- Ready thread with highest priority gets the processor
- When a thread is added to the ready list that has a higher priority than the currently running thread, immediately yield the processor to the new thread
- When threads are waiting for a lock, semaphore or a condition variable, the highest priority waiting thread should be woken up first
- Implementation details
  - compare priority of the thread being added to the ready list with that of the running thread
  - select next thread to run based on priorities
  - compare priorities of waiting threads when releasing locks, semaphores, condition variables

# Priority Based Scheduling



# thread\_yield() to implement preemption



- Current thread ("RUNNING") is moved to READY state, added to READY list.
- Then scheduler is invoked. Picks a new READY thread from READY list.
- Case a): there's only 1 READY thread. Thread is rescheduled right away
- Case b): there are other READY thread(s)
  - b.1) another thread has higher priority – it is scheduled
  - b.2) another thread has same priority – it is scheduled provided the previously running thread was inserted in tail of ready list.
- "thread\_yield()" is a call you can use whenever you identify a need to preempt current thread.
- **Exception:** inside an interrupt handler, use "intr\_yield\_on\_return()" instead

# Priority Inversion

- Strict priority scheduling can lead to a phenomenon called “priority inversion”
- Supplemental reading:
  - What really happened to the Pathfinder on Mars?
- Consider the following example where  $\text{prio}(H) > \text{prio}(M) > \text{prio}(L)$ 
  - H needs a lock currently held by L, so H blocks
  - M that was already on the ready list gets the processor before L
  - H indirectly waits for M
  - (on Path Finder, a watchdog timer noticed that H failed to run for some time, and continuously reset the system)



## Task 2B: Priority Donation

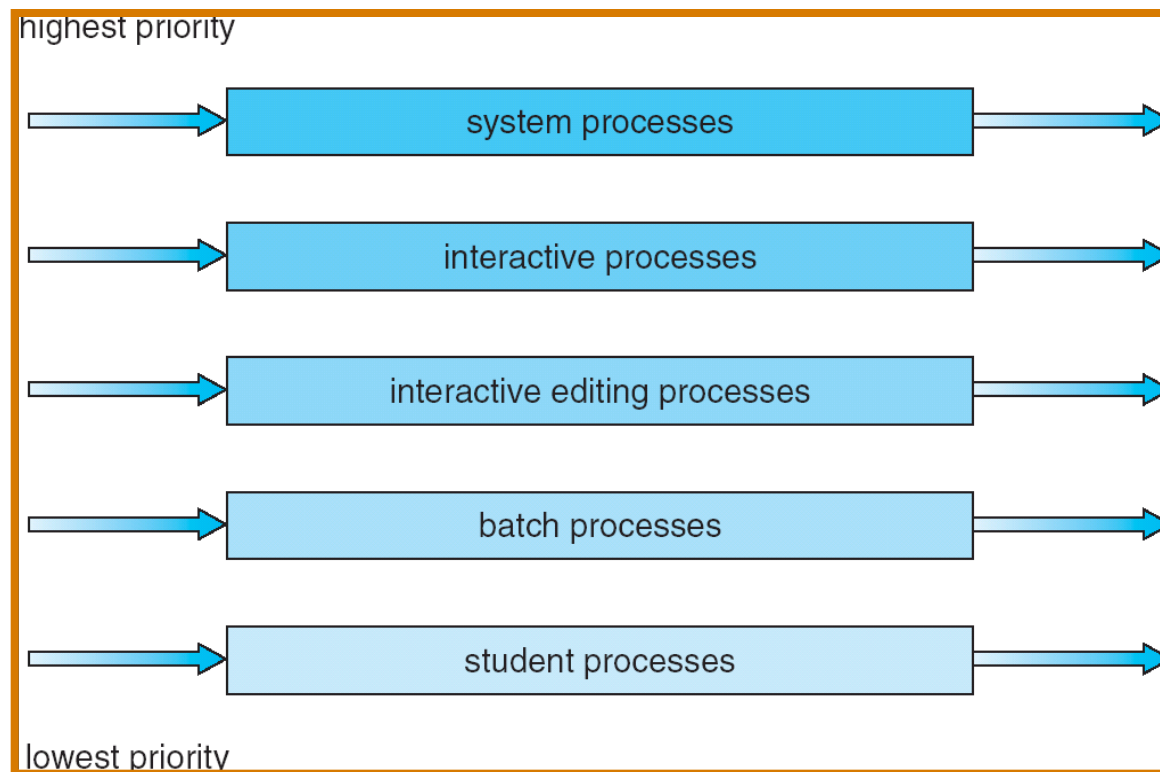
- When a high priority thread H waits on a lock held by a lower priority thread L, donate H's priority to L and recall the donation once L releases the lock
- Implement priority donation for locks
- Important:
  - Remember to return L to previous priority once it releases the lock.
  - Be sure to handle multiple donations (max of all donations)
  - Be sure to handle nested donations, e.g., H waits on M which waits on L...



## Task 3: Advanced Scheduler

- Implement Multi Level Feedback Queue Scheduler
- Priority donation not needed in the advanced scheduler – two implementations are not required to coexist
  - Only one is active at a time
- Advanced Scheduler must be chosen only if ‘-mlfqs’ kernel option is specified
- Read section on 4.4 BSD Scheduler in the Pintos manual for detailed information
- Some of the parameters are real numbers and calculations involving them have to be simulated using integers.
  - Write a fixed-point layer (header file)

# Multilevel Queue Scheduling




# Multilevel Feedback Queue

- A process can move between the various queues; aging can be implemented this way
- Multilevel-feedback-queue scheduler defined by the following parameters:
  - number of queues
  - scheduling algorithms for each queue
  - method used to determine which queue a process will enter when that process needs service
  - method used to determine when to upgrade a process
  - method used to determine when to degrade a process

# Example of Multilevel Feedback Queue

- Three queues:
  - $Q_0$  - RR with  $q = 8$  ms
  - $Q_1$  - RR with  $q = 16$  ms
  - $Q_2$  - FCFS



- Scheduling
  - A new job enters queue  $Q_0$  which is served FCFS. When it gains CPU, job receives 8 milliseconds. If it does not finish in 8 milliseconds, job is moved to queue  $Q_1$ .
  - At  $Q_1$  job is again served FCFS and receives 16 additional milliseconds. If it still does not complete, it is preempted and moved to queue  $Q_2$ .

# Suggested Order of Implementation

- Alarm Clock
  - easier to implement compared to the other parts
  - other parts not dependent on this
- Priority Scheduler
  - needed for implementing Priority Donation and Advanced Scheduler
- Priority Donation | Advanced Scheduler
  - these two parts are independent of each other
  - can be implemented in any order but only after Priority Scheduler is ready

# Debugging your Code

- printf, ASSERT, backtraces, gdb
- Running pintos under gdb
  - Invoke pintos with the gdb option  
`pintos --gdb -- run testname`
  - On another terminal invoke gdb  
`pintos-gdb kernel.o`
  - Issue the command  
`debugpintos`
  - All the usual gdb commands can be used: step, next, print, continue, break, clear etc
  - Use the pintos debugging macros described in manual

## Step 4: Testing

Pintos provides a very systematic testing suite for your project:

1. Run all tests:

```
$ make check
```

2. Run individual tests:

```
$ make tests/threads/alarm-multiple.result
```

3. Run the grading script:

```
$ make grade
```

# make check

- **pass** tests/threads/alarm-single
- **pass** tests/threads/alarm-multiple
- **pass** tests/threads/alarm-simultaneous
- FAIL tests/threads/alarm-priority
- **pass** tests/threads/alarm-zero
- **pass** tests/threads/alarm-negative
- FAIL tests/threads/priority-change
- FAIL tests/threads/priority-donate-one
- FAIL tests/threads/priority-donate-multiple
- FAIL tests/threads/priority-donate-multiple2
- FAIL tests/threads/priority-donate-nest
- FAIL tests/threads/priority-donate-sema
- FAIL tests/threads/priority-donate-lower
- FAIL tests/threads/priority-fifo
- FAIL tests/threads/priority-preempt
- FAIL tests/threads/priority-sema
- FAIL tests/threads/priority-condvar
- FAIL tests/threads/priority-donate-chain
- FAIL tests/threads/mlfqs-load-1
- FAIL tests/threads/mlfqs-load-60
- FAIL tests/threads/mlfqs-load-avg
- FAIL tests/threads/mlfqs-recent-1
- **pass** tests/threads/mlfqs-fair-2
- **pass** tests/threads/mlfqs-fair-20
- FAIL tests/threads/mlfqs-nice-2
- FAIL tests/threads/mlfqs-nice-10
- FAIL tests/threads/mlfqs-block



# Grading

**TOTAL 110 points:** 93 points for the implementation + 17 points for the documentation:

- **18 points:** A completely working Alarm Clock implementation that passes all six (6) tests.
- **38 points:** A fully functional Priority Scheduler that passes all twelve (12) tests.
- **37 points:** A working advanced scheduler that passes all nine (9) tests.
- **12 points:** A complete design document.
- **5 points:** A well-documented and clean source code.

# make grade (1)

## SUMMARY OF INDIVIDUAL TESTS

Functionality and robustness of alarm clock (tests/threads/Rubric.alarm):

4/ 4 tests/threads/alarm-single  
4/ 4 tests/threads/alarm-multiple  
4/ 4 tests/threads/alarm-simultaneous  
4/ 4 tests/threads/alarm-priority  
  
1/ 1 tests/threads/alarm-zero  
1/ 1 tests/threads/alarm-negative

- Section summary.

6/ 6 tests passed  
18/ 18 points subtotal

Functionality of priority scheduler (tests/threads/Rubric.priority):

3/ 3 tests/threads/priority-change  
3/ 3 tests/threads/priority-preempt  
  
3/ 3 tests/threads/priority-fifo  
3/ 3 tests/threads/priority-sema  
3/ 3 tests/threads/priority-condvar

## make grade (2)

TOTAL TESTING SCORE: 100.0%  
ALL TESTED PASSED -- PERFECT SCORE

---

### SUMMARY BY TEST SET

Test Set	Pts	Max	% Ttl	% Max
tests/threads/Rubric.alarm	18/	18	20.0%/	20.0%
tests/threads/Rubric.priority	38/	38	40.0%/	40.0%
tests/threads/Rubric.mlfqs	37/	37	40.0%/	40.0%
Total			100.0%/	100.0%

---

*Pintos include fully automated grading scripts, students see score before submission*

# Step 5: Design Document

Use the template in Section 2.2.1 of the Pintos documentation.

```
ALARM CLOCK
=====

---- DATA STRUCTURES ----

>> A1: Copy here the declaration of each new or changed `struct' or
>> `struct' member, global or static variable, `typedef', or
>> enumeration. Identify the purpose of each in 25 words or less.

---- ALGORITHMS ----

>> A2: Briefly describe what happens in a call to timer_sleep(),
>> including the effects of the timer interrupt handler.

>> A3: What steps are taken to minimize the amount of time spent in
>> the timer interrupt handler?

---- SYNCHRONIZATION ----

>> A4: How are race conditions avoided when multiple threads call
>> timer_sleep() simultaneously?

>> A5: How are race conditions avoided when a timer interrupt occurs
>> during a call to timer_sleep()?

---- RATIONALE ----

>> A6: Why did you choose this design? In what ways is it superior to
>> another design you considered?
```

# Submission

1. All source code (the full source tree)
  2. README file
  3. Design document
- Everything **due by October 22nd @11:59PM**