

CSE 421/521 - Operating Systems
Fall 2014

LECTURE - XVII

VIRTUAL MEMORY - II

Tevfik Koşar

University at Buffalo
October 23rd, 2014

Roadmap

- Virtual Memory
 - Performance of Demand Paging
 - Page Replacement Algorithms
 - Least Recently Used (LRU)
 - LRU Approximations
 - Counting Algorithms
 - Allocation Policies
 - Thrashing
 - Working Set Model



Performance of Demand Paging

- Page Fault Rate $0 \leq p \leq 1.0$
 - if $p = 0$ no page faults
 - if $p = 1$, every reference is a fault

- Effective Access Time (EAT)

$$\begin{aligned} \text{EAT} = & (1 - p) \times \text{memory access} \\ & + p \times (\text{page fault overhead} \\ & \quad + [\text{swap page out}] \\ & \quad + \text{swap page in} \\ & \quad + \text{restart overhead}) \end{aligned}$$

Demand Paging Example

- Memory access time = 1 microsecond
- 50% of the time the page that is being replaced has been modified and therefore needs to be swapped out
- Swap Page Time = 10 msec = 10,000 microsec
- EAT = ?

Demand Paging Example

- Memory access time = 1 microsecond
- 50% of the time the page that is being replaced has been modified and therefore needs to be swapped out
- Swap Page Time = 10 msec = 10,000 microsec
- $$\begin{aligned} \text{EAT} &= (1 - p) \times 1 + p \times (10,000 + 1/2 \times 10,000) \\ &= 1 + 14,999 \times p \quad (\text{in microsec}) \end{aligned}$$
- What if 1 out of 1000 memory accesses cause a page fault?
- What if we only want 30% performance degradation?

FIFO

- FIFO is obvious, and simple to implement
 - when you page in something, put it on the tail of a list
 - evict page at the head of the list
- Why might this be good?
 - maybe the one brought in longest ago is not being used
- Why might this be bad?
 - then again, maybe it *is* being used
 - have absolutely no information either way
- In fact, FIFO's performance is typically lousy
- In addition, FIFO suffers from **Belady's Anomaly**
 - there are **reference strings** for which the fault rate *increases* when the process is given more physical memory

Optimal (Belady's) Algorithm

- **Provably optimal:** lowest fault rate (remember SJF?)
 - evict the page that won't be used for the longest time in future
 - **problem: impossible to predict the future**
- Why is Belady's Optimal algorithm useful?
 - as a yardstick to compare other algorithms to optimal
 - if Belady's isn't much better than yours, yours is pretty good
 - how could you do this comparison?
- Is there a best practical algorithm?
 - no; depends on workload
- Is there a worst algorithm?
 - no, but random replacement does pretty badly
 - there are some other situations where OS's use near-random algorithms quite effectively!

Least Recently Used (LRU)

- Replace the page that has not been used for the longest amount of time in the past
- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5



Least Recently Used (LRU)

- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- 4 frames example:

1	5	
2		
3	5	4
4	3	

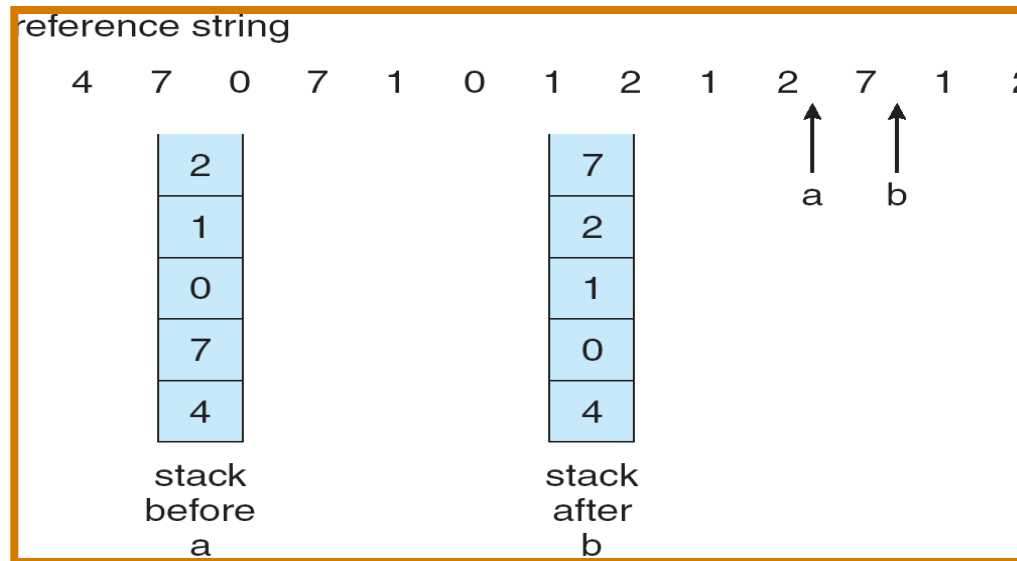
8 page faults

Least Recently Used (LRU)

- LRU uses reference information to make a more informed replacement decision
 - idea: past experience gives us a guess of future behavior
 - on replacement, evict the page that hasn't been used for the longest amount of time
 - LRU looks at the past, Belady's wants to look at future
 - *How is LRU different from FIFO?*
- Implementation
 - to be perfect, must grab a timestamp on every memory reference, then order or search based on the timestamps ...
 - way too costly in memory bandwidth, algorithm execution time, etc.
 - so, we need a cheap approximation ...

LRU Implementations

- **Stack implementation** - keep a stack of page numbers in a double link form:
 - Page referenced:
 - move it to the top
 - requires 6 pointers to be changed
 - No search for replacement



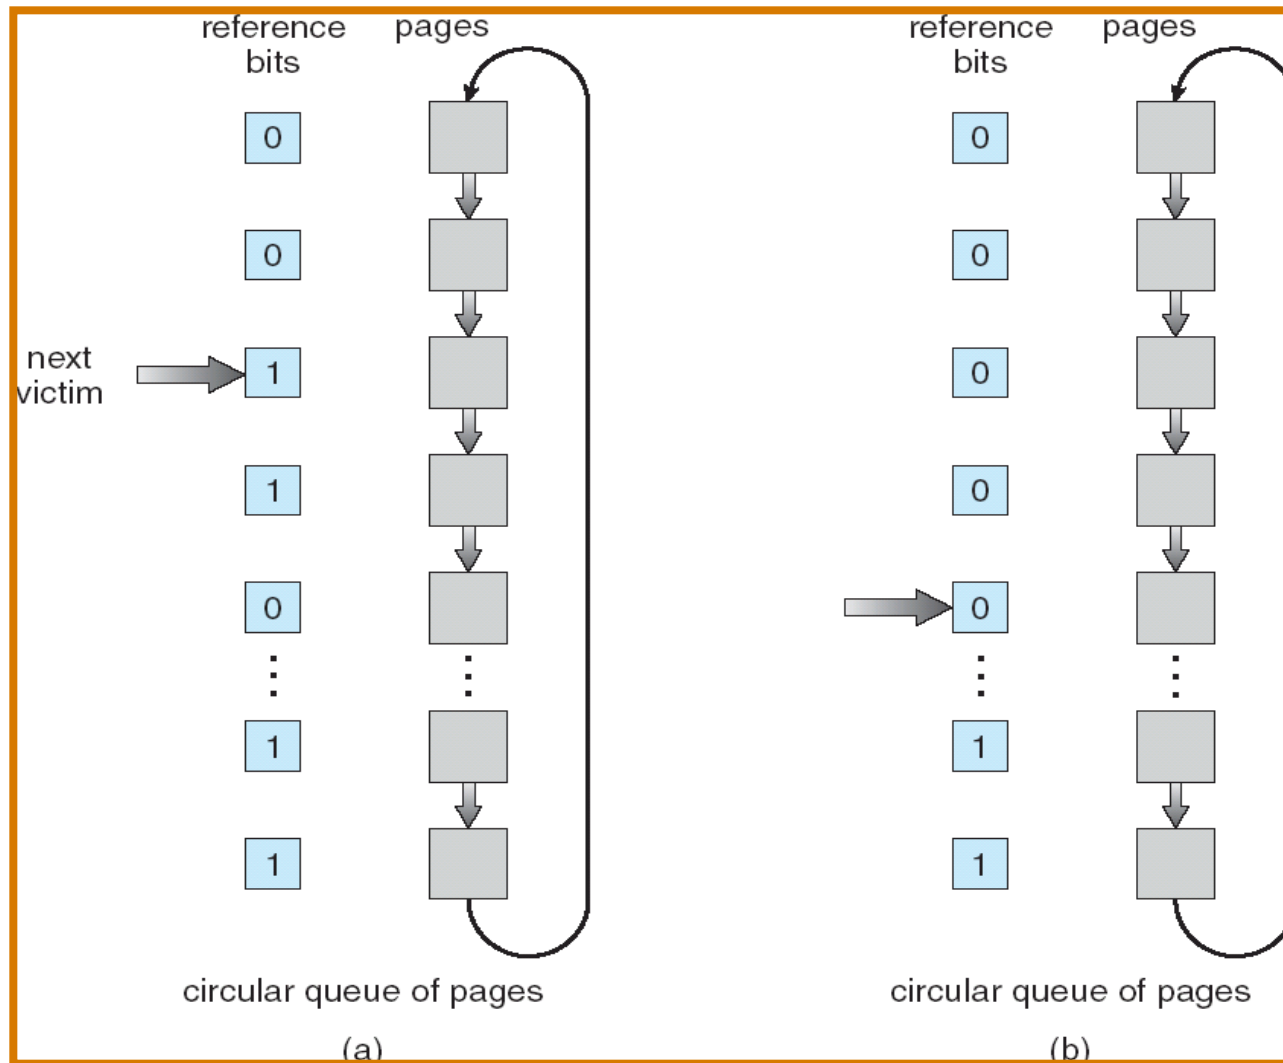
LRU Approximation Algorithms

- **Reference bit**
 - With each page associate a bit, initially = 0
 - When page is referenced bit set to 1
 - Replace the one which is 0 (if one exists). We do not know the order, however.
- **Additional Reference bits**
 - 1 byte for each page: eg. 00110011
 - Shift right at each time interval

LRU Clock Algorithm

- AKA Not Recently Used (NRU) or **Second Chance**
 - replace page that is “old enough”
 - logically, arrange all physical page frames in a big circle (clock)
 - just a circular linked list
 - a “clock hand” is used to select a good LRU candidate
 - sweep through the pages in circular order like a clock
 - if ref bit is off, it hasn’t been used recently, we have a victim
 - so, what is minimum “age” if ref bit is off?
 - if the ref bit is on, turn it off and go to next page
 - arm moves quickly when pages are needed
 - low overhead if have plenty of memory
 - if memory is large, “accuracy” of information degrades
 - add more hands to fix

Second-Chance (clock) Page-Replacement Algorithm



Counting Algorithms

- Keep a counter of the number of references that have been made to each page
- **LFU Algorithm**: replaces page with smallest count
- **MFU Algorithm**: based on the argument that the page with the smallest count was probably just brought in and has yet to be used

Exercise

- Consider the following page-reference string:

1, 2, 3, 4, 4, 3, 2, 1, 5, 6, 2, 1, 2, 3, 7, 8, 3, 2, 1, 5

Assuming 4 memory frames and LR, LFU, or Optimal page replacement algorithms, how many page faults, page hits, and page replacements would occur? Show your page assignments to frames.

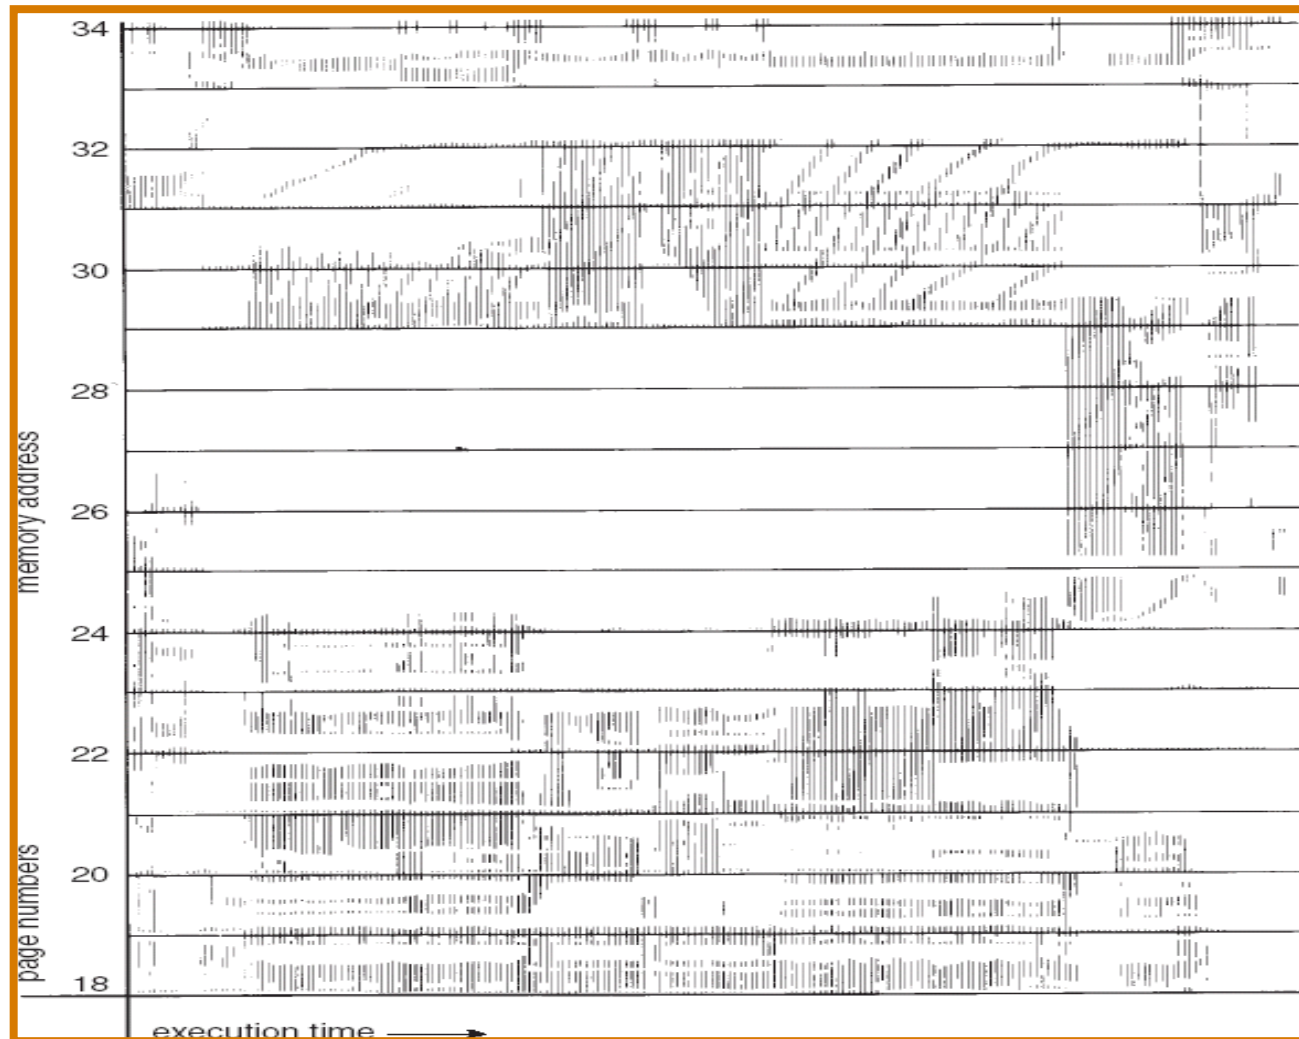
1	2	3	4	4	3	2	1	5	6	2	1	2	3	7	8	3	2	1	5

of page faults:

of page hits:

of page replacements:


Locality in a Memory-Reference Pattern



Working-Set Model

- $\Delta \equiv$ working-set window \equiv a fixed number of page references
Example: 10,000 instruction
- WSS_i (working set of Process P_i) =
total number of pages referenced in the most recent Δ (varies in time)
 - if Δ too small will not encompass entire locality
 - if Δ too large will encompass several localities
 - if $\Delta = \infty \Rightarrow$ will encompass entire program
- $D = \sum WSS_i \equiv$ total demand frames
- if $D > m \Rightarrow$ Thrashing
- Policy if $D > m$, then suspend one of the processes


Working-set model



Thrashing

- If a process does not have “enough” frames, the page-fault rate is very high. This leads to:
 - Replacement of active pages which will be needed soon again
 - **Thrashing** \equiv a process is busy swapping pages in and out
- Which will in turn cause:
 - low CPU utilization
 - operating system thinks that it needs to increase the degree of multiprogramming
 - another process added to the system

Thrashing (Cont.)



Allocation of Frames

- Each process needs *minimum* number of pages
- Two major allocation schemes
 - fixed allocation
 - priority allocation

Fixed Allocation

- **Equal allocation** - For example, if there are 100 frames and 5 processes, give each process 20 frames.
- **Proportional allocation** - Allocate according to the size of process

s_i = size of process p_i

$$S = \sum s_i$$

m = total number of frames

$$a_i = \text{allocation for } p_i = \frac{s_i}{S} \times m$$

$$m = 64$$

$$s_1 = 10$$

$$s_2 = 127$$

$$a_1 = \frac{10}{137} \times 64 \approx 5$$

$$a_2 = \frac{127}{137} \times 64 \approx 59$$

Priority Allocation


- Use a proportional allocation scheme using priorities rather than size
- If process P_i generates a page fault,
 - select for replacement one of its frames
 - select for replacement a frame from a process with lower priority number

Global vs. Local Allocation

- **Global replacement** - process selects a replacement frame from the set of all frames; one process can take a frame from another
- **Local replacement** - each process selects from only its own set of allocated frames

Summary

- Virtual Memory
 - Page Replacement Algorithms
 - Optimal Algorithm
 - Least Recently Used (LRU)
 - LRU Approximations
 - Counting Algorithms
 - Allocation Policies
 - Thrashing
 - Working Set Model



- Next Lecture: Project 2 Discussion
- Reading Assignment: Chapter 9 from Silberschatz.

Acknowledgements

- “Operating Systems Concepts” book and supplementary material by A. Silberschatz, P. Galvin and G. Gagne
- “Operating Systems: Internals and Design Principles” book and supplementary material by W. Stallings
- “Modern Operating Systems” book and supplementary material by A. Tanenbaum
- R. Doursat and M. Yuksel from UNR
- Gribble, Lazowska, Levy, and Zahorjan from UW