

CSE 421/521 - Operating Systems
Fall 2014 Recitations

RECITATION - I
UNIX C PROGRAMMING
PROF. TEVFIK KOSAR

Presented by Sharath Chandrashekhara

University at Buffalo
September 2-3, 2014

logon

- `ssh timberlake.cse.buffalo.edu -l username`
 - or:
- `ssh username@timberlake.cse.buffalo.edu`
- `passwd`: change password
- `putty`: a free telnet/ssh client
- `ls /bin` (`ls /usr/bin`)
- `man ...`
- text editing: `vi`, `emacs`, `pico`

Vi Editor

- vi filename
 - a: enter *insert* mode, after the cursor
 - i: enter *insert* mode, before the cursor
 - O: enter *insert* mode, above the cursor
 - o: enter *insert* mode, below the cursor
 - r: replace one character under the cursor
 - u: undo the last change to the file.
 - x: delete character under the cursor
 - yy: copy line
 - dd: delete line
 - :w: write
 - :q: quit
 - :q!: quit without saving changes
 - /keyword : search for the keyword in text
 - :n : go to line number n
- Vi tutorial: <http://www.gnulamp.com/vi.html>

Emacs Editor

- Emacs filename
 - CTRL-d : delete one character
 - CTRL-k : delete one line
 - CTRL-y : paste
 - CTRL-x 2 : split window into 2 (horizontal)
 - CTRL-x 3 : split window into 2 (vertical)
 - CTRL-x o : switch window
 - CTRL-x 1 : kill all other windows
 - CTRL-x u : undo (also CTRL-_)
 - CTRL-x CTRL-f: open file
 - CTRL-x CTRL-b: open buffer (CTRL-x b: switch to buffer)
 - CTRL-s : search
 - CTRL-x CTRL-s: save file
 - CTRL-x CTRL-c: quit
- Emacs Tutorial: http://www.gnu.org/software/emacs/tour/emacs_toc.html

Or...

- Use any editor you are familiar with.
(Notepad, Wordpad, etc.)
- After file is written, upload the file using SFTP software such as FileZilla

Files and Directories

- directory operations
 - ls: list
 - cd: change directory
 - pwd: print working directory
 - mkdir: create directory
 - rmdir: remove directory
- file operations
 - cp: copy
 - rm: delete
 - mv: move (rename)
 - cat, more, less: examine
- file permissions: rwx rwx rwx
 user group others
 - chmod 755 filename (or chmod u+r filename) (or chmod u=rwx)

Processes

- ps : list currently active user processes
- ps aux: list all active processes in long format
- kill n : kill process with id=n
- kill -9 n : force to kill

- CTRL-z : push to background
- fg : bring to foreground (also fg n: bring nth process)

- top: system utilization information
- time command : calculate time for a given command

Basic C Program: Print to stdout

```
#include <stdio.h>
```

```
main()
```

```
{
```

```
    printf("Hello, CSC4304 Class!\n");
```

```
}
```

```
---
```

```
gcc prog1.c          ==> a.out
```

```
gcc prog1.c -o prog1 ==> prog1
```

```
make prog1           ==> prog1
```


Header Files

- The C compiler works in 3 phases:
 - ① Pre-process source files
 - ② Compile source files into object files
 - ③ Link object files into an executable
- `#include <stdio.h>` means "include the contents of standard file `stdio.h` here"
 - ① Standard files are usually located in directory `/usr/include`
 - ② `/usr/include/stdio.h` may contain `#include` statements itself...
- You can use `#include` to include your own files into each other:
 - ▶ `#include "myfile.h"` means: "include file `myfile.h` (from the current directory) here"
 - ▶ Included files usually have extension `".h"` (header)

Basic Data Types

- Basic Types
 - char : character - 1 byte
 - short: short integer - 2 bytes
 - int: integer - 4 bytes
 - long: long integer - 4 bytes
 - float: floating point - 4 bytes
 - double - double precision floating point - 8 bytes
- Formatting Template
 - %d: integers
 - %f: floating point
 - %c: characters
 - %s: string
 - %x: hexadecimal
 - %u: unsigned int

Test Size of Data Types

```
#include <stdio.h>
```

```
main()
```

```
{
```

```
    printf("sizeof(char): %d\n", sizeof(char));
```

```
    printf("sizeof(short): %d\n", sizeof(short));
```

```
    printf("sizeof(int): %d\n", sizeof(int));
```

```
    printf("sizeof(long): %d\n", sizeof(long));
```

```
    printf("sizeof(float): %d\n", sizeof(float));
```

```
    printf("sizeof(double): %d\n", sizeof(double));
```

```
}
```

Formatting

```
#include <stdio.h>
```

```
main()
```

```
{
```

```
    char var1;
```

```
    float f;
```

```
    printf(" Enter a character:");
```

```
    scanf("%c", &var1);
```

```
    printf("You have entered character:%c \n ASCII value=%d \n  
        Address=%x\n", var1, var1, &var1);
```

```
    printf(" And its float value would be: %.2f\n", (float)var1);
```

```
}
```

1
2

Formatting (*cont.*)

<pre>#include <stdio.h> int main(void) { int val = 5; char c = 'a'; char str[] = "world"; printf("Hello world\n"); printf("Hello %d World\n", val); printf("%d %c World\n", val, c); printf("Hello %s\n", str); printf("Hello %d\n", str); return 0; }</pre>		
		Hello world
		Hello 5 World
		5 a World
		Hello world
		** wrong! **

Read argument and print

```
#include <stdio.h>

main(int argc, char* argv[])
{
    if (argc < 2){
        printf("Usage: %s <your name>\n", argv[0]);
    }
    else{
        printf("Hello, %s!\n", argv[1]);
    }
}
```

Read from stdin and print

```
#include <stdio.h>
```

```
main()
```

```
{
```

```
    char name[64];
```

```
    printf("What's your name?");
```

```
    scanf("%s", name);
```

```
    printf("Hello, %s!\n", name);
```

```
}
```

Arrays

- Defining an array is easy:

```
int a[3];    /* a is an array of 3 integers */
```

- Array indexes go from 0 to n-1:

```
a[0] = 2; a[1] = 4; a[2] = a[0] + a[1];  
int x = a[a[0]];    /* what is the value of x? */
```

- ▶ **Beware:** in this example a[3] does not exist, but your compiler will not complain if you use it!

- ★ But your program may have a very strange behavior...

- You can create multidimensional arrays:

```
int matrix[3][2];  
matrix[0][1] = 42;
```


Strings

- A string is an array of characters:

```
char hello[15]="Hello, world!\n";
```

- Unlike in Java, you must decide in advance how many characters can be stored in a string.
 - ▶ You cannot change the size of the array afterwards
- Beware: strings are always terminated by a NULL character: '\0'
 - ▶ For example, "Hello" is string of 6 characters:

H	e	l	l	o	\0
---	---	---	---	---	----

Manipulating Arrays


- You cannot copy an array into another directly
 - ▶ You must copy each element one at a time

```
int a[3] = {12,24,36};  
int b[3];  
  
b = a;    /* This will NOT work! */  
  
b[0]=a[0];  
b[1]=a[1];  
b[2]=a[2]; /* This will work */
```

Manipulating Strings

- There are standard function to manipulate strings:
 - ▶ `strcpy(destination, source)` will copy string **source** into string **destination**:

```
char a[15] = "Hello, world!\n";  
char b[15];  
strcpy(b,a);
```

 Attention: `strcpy` does **not** check that **destination** is large enough to accomodate **source**.

```
char c[10];  
strcpy(c,a); /* This will get you in BIG trouble */
```

Manipulating Strings (*cont.*)

- Instead of `strcpy` it is always better to use `strncpy`:
 - ▶ `strncpy` takes one more parameter to indicate the maximum number of characters to copy:

```
char a[15] = "Hello, world!";  
char c[10];  
strncpy(c,a,9);    /* Why 9 instead of 10? */
```

Comparison Operators

- The following operators are defined for basic data types:

```
if (a == b) { ... }  
if (a != b) { ... }  
if (a < b) { ... }  
if (a <= b) { ... }  
if (a > b) { ... }  
if (a >= b) { ... }  
if ((a==b) && (c>d)) {...} /* logical AND */  
if ((a==b) || (c>d)) {...} /* logical OR */
```

- There is no boolean type in C. We use integers instead:
 - ▶ 0 means FALSE
 - ▶ Any other value means TRUE

```
int x;  
if (x) {...}           /* Equivalent to: if (x!=0) {...} */  
if (!x) {...}          /* Equivalent to: if (x==0) {...} */
```

Example

```
#include <stdio.h>
```

```
main()
```

```
{
```

```
    int x = 5;
```

```
    int y = 3;
```

```
    if (x=y){
```

```
        printf("x is equal to y, x=%d, y=%d\n", x, y);
```

```
    }
```

```
    else{
```

```
        printf("x is not equal to y, x=%d, y=%d\n", x, y);
```

```
    }
```

```
}
```

Classical Bugs

- Do not confuse '=' and '=='!

```
if (x=y) { ... }    /* This is correct C but it means something different */  
if (x=3) { /* always executed */ }  
if (x=0) { /* never executed */ }
```

- Do not confuse '&' and '&&'!

```
if (x&y) { ... }    /* This is correct C but it means something different */  
if (x|y) { ... }
```

Exercise:

- (7 & 8) vs (7 && 8)
- (7 | 8) vs (7 || 8)

Loops

```
while (x>0){
```

```
...
```

```
}
```

```
do{
```

```
...
```

```
} while (x>0);
```

```
for (x=0; X<3;x++) {...}
```


Functions

- In C, functions can be defined in two ways:

```
int foo() {                                /* function foo returns an int */
    ...
    return 123;
}

void bar(int p1, double p2) {              /* function bar returns nothing */
    ...
}
```

- Calling a function is easy:

```
int i = foo(); /* call function foo() */
bar(2, -4.321); /* call function bar() */
```

Memory Manipulation in C

- To a C program, memory is just a row of bytes
- Each byte has some value, and an address in the memory


123	2	45	254	2	66	67	234	99	1	0	0	12	92	15
20	21	22	23	24	25	26	27	28	29	30	31	32	33	34

Memory Manipulation in C

- When you define variables:

```
int count;  
unsigned char c;
```

- Memory is reserved to store the variables
- And the compiler 'remembers their location'



Memory Manipulation in C

- As a result, each variable has two properties:

- ① The '**value**' stored in the variable

- 💡 If you use the name of the variable, you refer to the variable's value

- ② The '**address**' of the memory used to store this value

- ★ Similar to a reference in Java (but not exactly the same)

- 💡 A variable that stores the address of another variable is called a **pointer**

- Pointers can be declared using the * character

```
int *ptr;           /* Pointer to an int */
unsigned char *ch;  /* Pointer to an unsigned char */
struct ComplexNumber *c; /* Pointer to a struct ComplexNumber */
int **pp;           /* Pointer to a pointer to an int */
void *v;            /* Pointer to anything (use with care!) */
```

Defining Pointers

- To use pointers, you must give them a value first
 - ▶ Like any other variable
- The '&' operator gives you the **memory address** of any variable

```
int i = 8;

int *p;      /* p is a pointer to an int */

p = &i;      /* p contains the address of variable i */

double *d = &i; /* ERROR, wrong pointer type */
```

Using Pointers

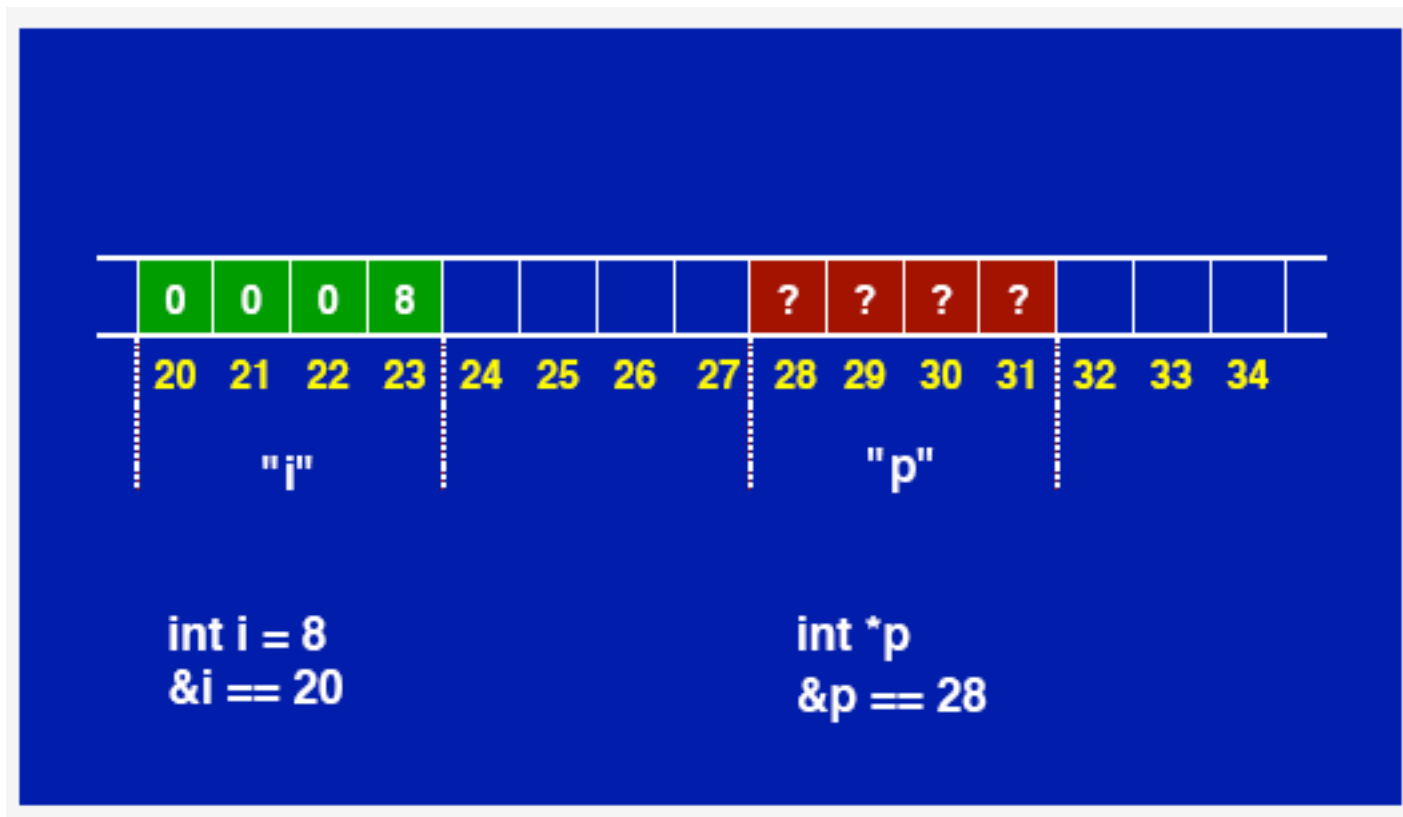
- Once you have a pointer, you can access the value of the variable being pointed by using '*'

```
int i = 8;  
int *p = &i;  
int j = *p;  
*p = 12;
```

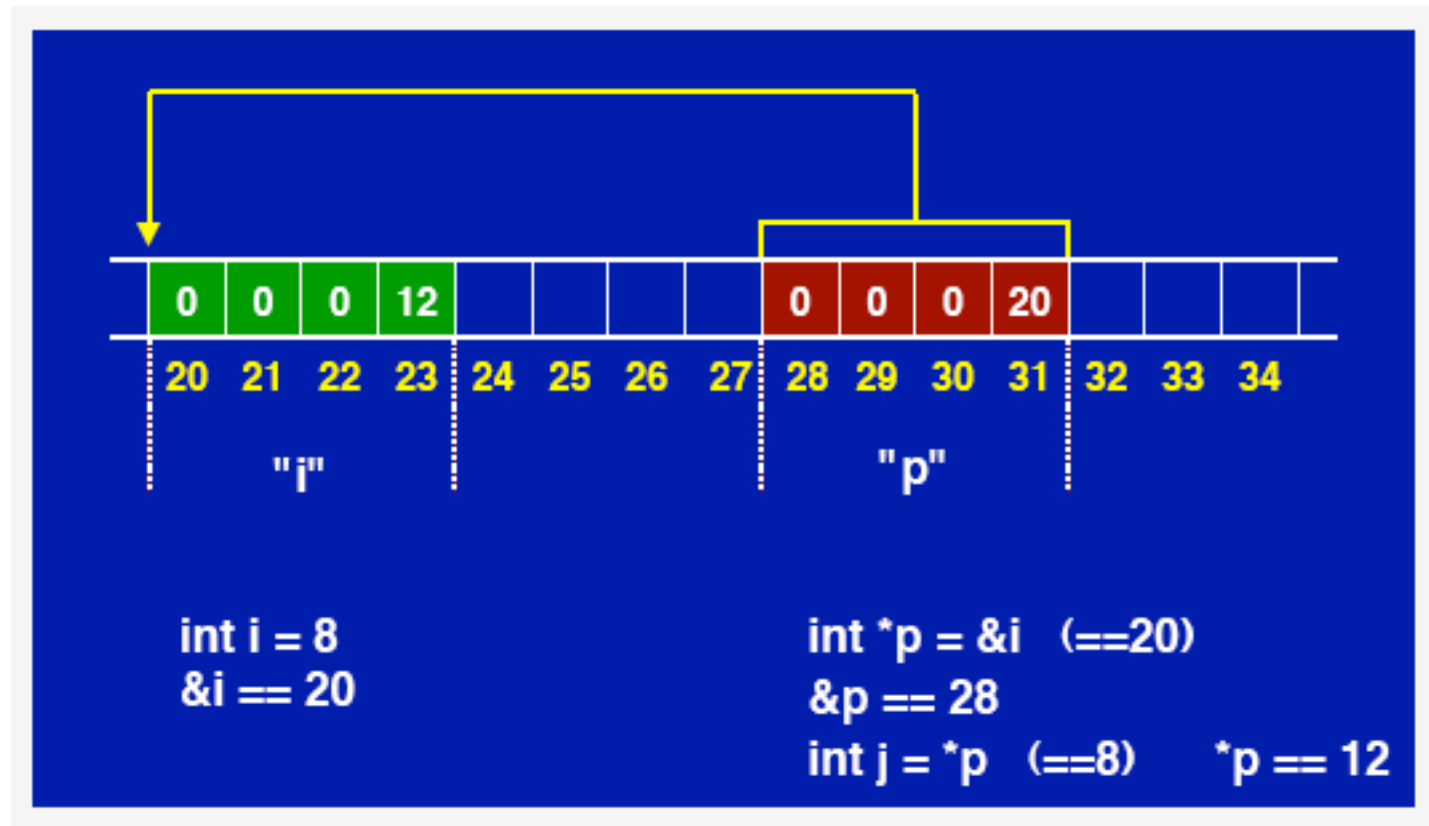
📌 Attention, the '*' sign is used for two different things:

- ▶ To **declare** a pointer variable: `int *p;`
- ▶ To **dereference** a pointer: `*p=12;`

Using Pointers



Using Pointers



Parameter Passing in C

- In C, function parameters are passed **by value**
 - ▶ Each parameter is copied
 - ▶ The function can access the copy, not the original value

```
#include <stdio.h>

void swap(int x, int y) {
    int temp = x;
    x = y;
    y = temp;
}

int main() {
    int x = 9;
    int y = 5;
    swap(x, y);
    printf("x=%d y=%d\n", x, y);
    return 0;
}
```

Parameter Passing in C

- In C, function parameters are passed **by value**
 - ▶ Each parameter is copied
 - ▶ The function can access the copy, not the original value

```
#include <stdio.h>

void swap(int x, int y) {
    int temp = x;
    x = y;
    y = temp;
}

int main() {
    int x = 9;
    int y = 5;
    swap(x, y);
    printf("x=%d y=%d\n", x, y); /* This will print: x=9 y=5 */
    return 0;
}
```

Parameter Passing in C

- To pass parameters by reference, use pointers
 - ▶ The pointer is copied
 - ▶ But the copy still points to the same memory address

```
#include <stdio.h>

void swap(int *x, int *y) {
    int temp = *x;
    *x = *y;
    *y = temp;
}

int main() {
    int x = 9;
    int y = 5;
    swap(&x, &y);
    printf("x=%d y=%d\n", x, y); /* This will print: x=5 y=9 */
    return 0;
}
```

Arrays and Pointers

- You can use pointers instead of arrays as parameters

```
#include <stdio.h>

void func1(int p[], int size) { }

void func2(int *p, int size) { }

int main() {
    int array[5];
    func1(array, 5);
    func2(array, 5);
    return 0;
}
```

Arrays and Pointers

- You can even use array-like indexing on pointers!

```
void clear(int *p, int size) {  
    int i;  
    for (i=0;i<size;i++) {  
        p[i] = 0;  
    }  
}  
  
int main() {  
    int array[5];  
    clear(array, 5);  
    return 0;  
}
```

Arrays and Pointers

- So a string is in fact just a pointer to a character array:

```
int main() {  
    char s1[32] = "Hello, world!\n";  
    char *s2;  
    char s3[32];  
    s2 = s1;           /* s1 and s2 point to the same character array */  
    strncpy(s3,s1,31); /* s3 contains a copy of s1 */  
}
```

Pointer Arithmetic


- Pointers are just a special kind of variable
- You can do **calculations** on pointers
 - ▶ You can use +, -, ++, -- on pointers
 - ▶ This has no equivalent in Java
- Be careful, operators work with the **size** of variable types!

```
int i = 8;  
int *p = &i;  
p++; /* increases p with sizeof(int) */  
  
char *c;  
c++; /* increases c with sizeof(char) */
```

Pointer Arithmetic


- This is obvious when using pointers as arrays:

```
int i;  
int array[5];  
int *p = array;  
  
for (i=0; i<5; i++) {  
    *p = 0;  
    p++;  
}
```



```
int array[5];  
int *p = array;
```


Pointer Arithmetic




```
int array[5];
```

```
int *p = array;
```

```
p++;
```

```
p++;
```



```
char array[5];
```

```
char *p = array;
```

```
p++;
```

```
p++;
```

Structures

- You can build higher-level data types by creating structures:

```
struct Complex {  
    float real;  
    float imag;  
};  
struct Complex number;  
number.real = 3.2;  
number.imag = -2;  
  
struct Parameter {  
    struct Complex number;  
    char description[32];  
};  
struct Parameter p;  
p.number.real = 42;  
p.number.imag = 12.3;  
strncpy(p.description, "My nice number", 31);
```

Pointers to Structures

- We very often use statements like:

```
(*pointer).field = value;
```

- There is another notation which means exactly the same:

```
pointer->field = value;
```

- For example:

```
struct data {  
    int counter;  
    double value;  
};  
  
void add(struct data *d, double value) {  
    d->counter++;  
    d->value += value;  
}
```

Enumerations

- **enum** is used to create a number of related constants

```
enum workdays {monday, tuesday, wednesday, thursday, friday };  
  
enum workdays today;  
today = tuesday;  
today = friday;  
  
enum weekend {saturday = 10, sunday = 20};
```

Variables

- C has two kinds of variables:
 - ▶ Local (declared inside of a function)
 - ▶ Global (declared outside of a function)

```
int global;  
  
void function() {  
    int local;  
}
```

Static Local Variables

- Declaring a static variable means it will persist across multiple calls to the function

```
void foo() {  
    static int i=0;  
    i++;  
    printf("i=%d\n",i); /* This prints the value of i on the screen */  
}  
  
int main() {  
    int i;  
    for (i=0;i<3;i++) foo();  
}
```

This program will output this:

```
i=1  
i=2  
i=3
```

Non-static Local Variables

- If i is not static, the same example program (from prev. slide) will output:
 - $i=1$
 - $i=1$
 - $i=1$

Global Variables

Global variables have file scope:

```
int i=0;
```

```
void foo() {  
    i++;  
    printf("i=%d\n",i);  
}
```

```
int main() {  
    for (i=0;i<3;i++) foo();  
}
```


Dynamic Memory Management

- Until now, all data have been static
 - ▶ It is clear by reading the program how much memory must be allocated
 - ▶ Memory is reserved at compile time
- But sometimes you want to specify the amount of memory to allocate **at runtime!**
 - ▶ You need a string, but you don't know yet how long it will be
 - ▶ You need an array but you don't know yet how many elements it should contain
 - ▶ Sizes depend on run-time results, user input, etc.

Dynamic Memory Management

- `malloc()` will allocate any amount of memory you want:

```
#include <stdlib.h>
void *malloc(size_t size);
```

- ▶ `malloc` takes a size (in bytes) as a parameter
 - ★ If you want to store 3 integers there, then you must reserve `3*sizeof(int)` bytes
- ▶ It returns a pointer to the newly allocated piece of memory
 - ★ It is of type `void *`, which means "pointer to anything"
 - ★ Do not store it as a `void *`! You should "cast" it into a usable pointer:

```
#include <stdlib.h>
int *i = (int *) malloc(3*sizeof(int));
i[0] = 12;
i[1] = 27;
i[2] = 42;
```

Dynamic Memory Management

- After you have used `malloc`, the memory will remain allocated until you decide to destroy it

```
#include <stdlib.h>
void free(void *pointer);
```

- After you have finished using dynamic memory, **you must release it!**
 - ▶ Otherwise it will remain allocated (and unused) until the end of the program's execution

```
int main() {
    int *i = (int *) malloc(3*sizeof(int));
    /* Use i */
    free(i);
    /* Do something else */
}
```

Dynamic Memory Management

- Unlike arrays, dynamically allocated memory can be returned from a function.

```
int *createIntArrayWrong() {  
    char tmp[32];  
    return tmp;                                /* WRONG! */  
}  
  
int *createIntArray(int size) {  
    return (int *) malloc(size*sizeof(int)); /* CORRECT */  
}  
  
int main() {  
    int *array = createIntArray(10);  
    /* ... */  
    free(array);  
    return 0;  
}
```

Memory Leaks

- You must **always** keep a pointer to allocated memory
 - ▶ You need this to use it, and free it later
 - ▶ If you don't, you've got a **memory leak**
 - ▶ Memory leaks will slowly reserve all the machine memory, causing the program (or the machine) to crash eventually!

```
int main() {
    int *i = (int *) malloc(3*sizeof(int));
    i = 0;      /* Woops, I lost the pointer to my dynamic memory */
    free(???); /* It is too late to free my dynamic memory */
}
```

- If you run out of memory, malloc will return NULL

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    int *array = (int *) malloc(10*sizeof(int));

    if (array == NULL) {
        printf("Out of memory!\n");
        return 1;
    }

    /* do something useful here */
    return 0;
}
```

malloc Example

```
int main ()
{
    int x = 11;
    int *p, *q;

    p = (int *) malloc(sizeof (int));
    *p = 66;
    q = p;
    printf ("%d %d %d\n", x, *p, *q);
    x = 77;
    *q = x + 11;
    printf ("%d %d %d\n", x, *p, *q);
    p = (int *) malloc(sizeof (int));
    *p = 99;
    printf ("%d %d %d\n", x, *p, *q);
}
```

\$./malloc
11 66 66
77 88 88
77 99 88

free Example

```
int main ()
{
    int x = 11;
    int *p, *q;
    p = (int *) malloc(sizeof (int));
    *p = 66;
    q = (int *) malloc(sizeof (int));
    *q = *p - 11;
    free(p);
    printf ("%d %d %d\n", x, *p, *q);
    x = 77;
    p = q;
    q = (int *) malloc(sizeof (int));
    *q = x + 11;
    printf ("%d %d %d\n", x, *p, *q);
    p = &x;
    p = (int *) malloc(sizeof (int));
    *p = 99;
    printf ("%d %d %d\n", x, *p, *q);
    q = p;
    free(q);
    printf ("%d %d %d\n", x, *p, *q);
}
```

./free		
11	?	55
77	55	88
77	99	88
77	?	?

Acknowledgments

- Advanced Programming in the Unix Environment by R. Stevens
- The C Programming Language by B. Kernighan and D. Ritchie
- Understanding Unix/Linux Programming by B. Molay
- Lecture notes from B. Molay (Harvard), T. Kuo (UT-Austin), G. Pierre (Vrije), M. Matthews (SC), and B. Knicki (WPI).